

Introduction to Computer Design

E. Sanchez

**Politecnico di Torino
Dipartimento di Automatica e Informatica**

Computer evolution

The first general-purpose computer was created in the late 40s.

A Personal Computer, that can now be bought for about \$500, is practically equivalent (in terms of performance and memory) to what could be bought for about \$1M in 1985.

This evolution has been made possible by

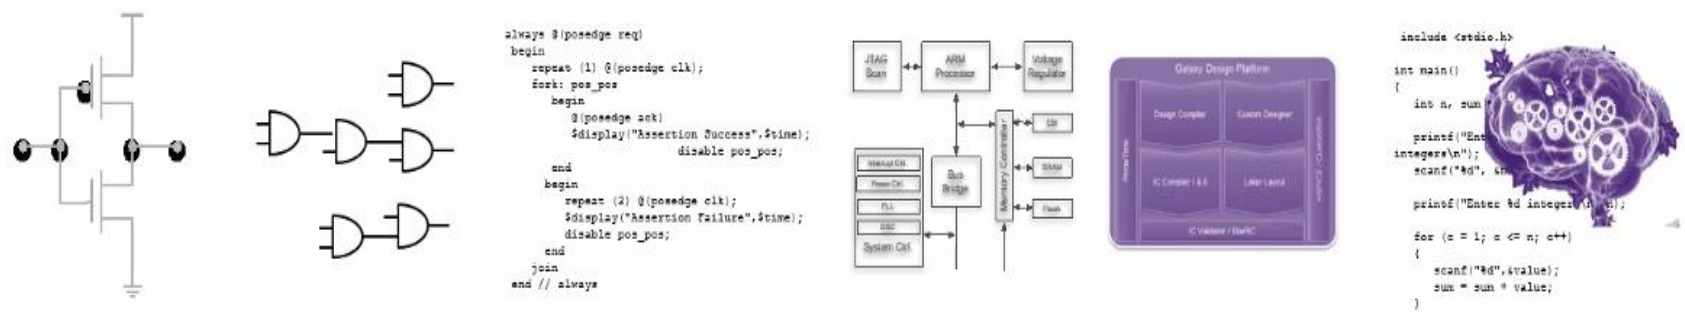
- **Advances in semiconductor technology**
- **Innovations in computer design**
- **Advances in software.**

Chip-design history



transistors 2,000 30,000 1,000,000 40,000,000 >10,000,000,000 HUGE

Design entry




Design complexity



Annual size of the global Datasphere

Figure 1 – Annual Size of the Global Datasphere




source: IDC Datasphere whitepaper


zettaByte ZB $\rightarrow 10^{21} \rightarrow 2^{70}$

Annual size of the global Datasphere

Figure 1 – Annual Size of the Global Datasphere



Microprocessor performance growth




7



Processor performance growth

The annual increase raised to more than 50% for the RISC architectures in the 1980s.




Starting from 2002, the yearly processor performance increase dropped to about 20% due to

- Intel Core i7 4 cores 4.2 GHz (Boost to 4.5 GHz)
Intel Core i7 4 cores 4.0 GHz (Boost to 4.2 GHz)
Intel Core i7 4 cores 4.0 GHz (Boost to 4.2 GHz)
Intel Xeon 4 cores 3.7 GHz (Boost to 4.1 GHz)
Intel Xeon 4 cores 3.6 GHz (Boost to 4.0 GHz)
Intel Xeon 4 cores 3.6 GHz (Boost to 4.0 GHz)
Intel Core i7 4 cores 3.4 GHz (boost to 3.8 GHz)
Intel Xeon 6 cores, 3.3 GHz (boost to 3.6 GHz)
Intel Xeon 4 cores, 3.3 GHz (boost to 3.6 GHz)
Intel Core i7 Extreme 4 cores 3.8 GHz (boost to 4.5 GHz)




Microprocessor performance

Since 2004, major industries canceled its race for high-performance single processor projects running at higher clock frequencies, and embrace multicore CPUs.




yth

Processor	Score
Intel Core i7 4 cores 4.2 GHz (Boost to 4.5 GHz)	49,935
Intel Core i7 4 cores 4.0 GHz (Boost to 4.2 GHz)	49,870
Intel Xeon 4 cores 3.6 GHz (Boost to 4.0 GHz)	39,935
Intel Xeon 4 cores 3.6 GHz (Boost to 4.0 GHz)	39,919
Intel Core i7 4 cores 3.4 GHz (boost to 3.8 GHz)	39,919
Intel Xeon 6 cores, 3.3 GHz (boost to 3.6 GHz)	39,919
Intel Xeon 4 cores, 3.3 GHz (boost to 3.6 GHz)	39,919
Intel Xeon 4 cores 3.7 GHz (Boost to 4.1 GHz)	31,999




Microprocessor performance growth


3.5% performance improvement, or doubling every 20 years. Is this the end of the Moore's law?



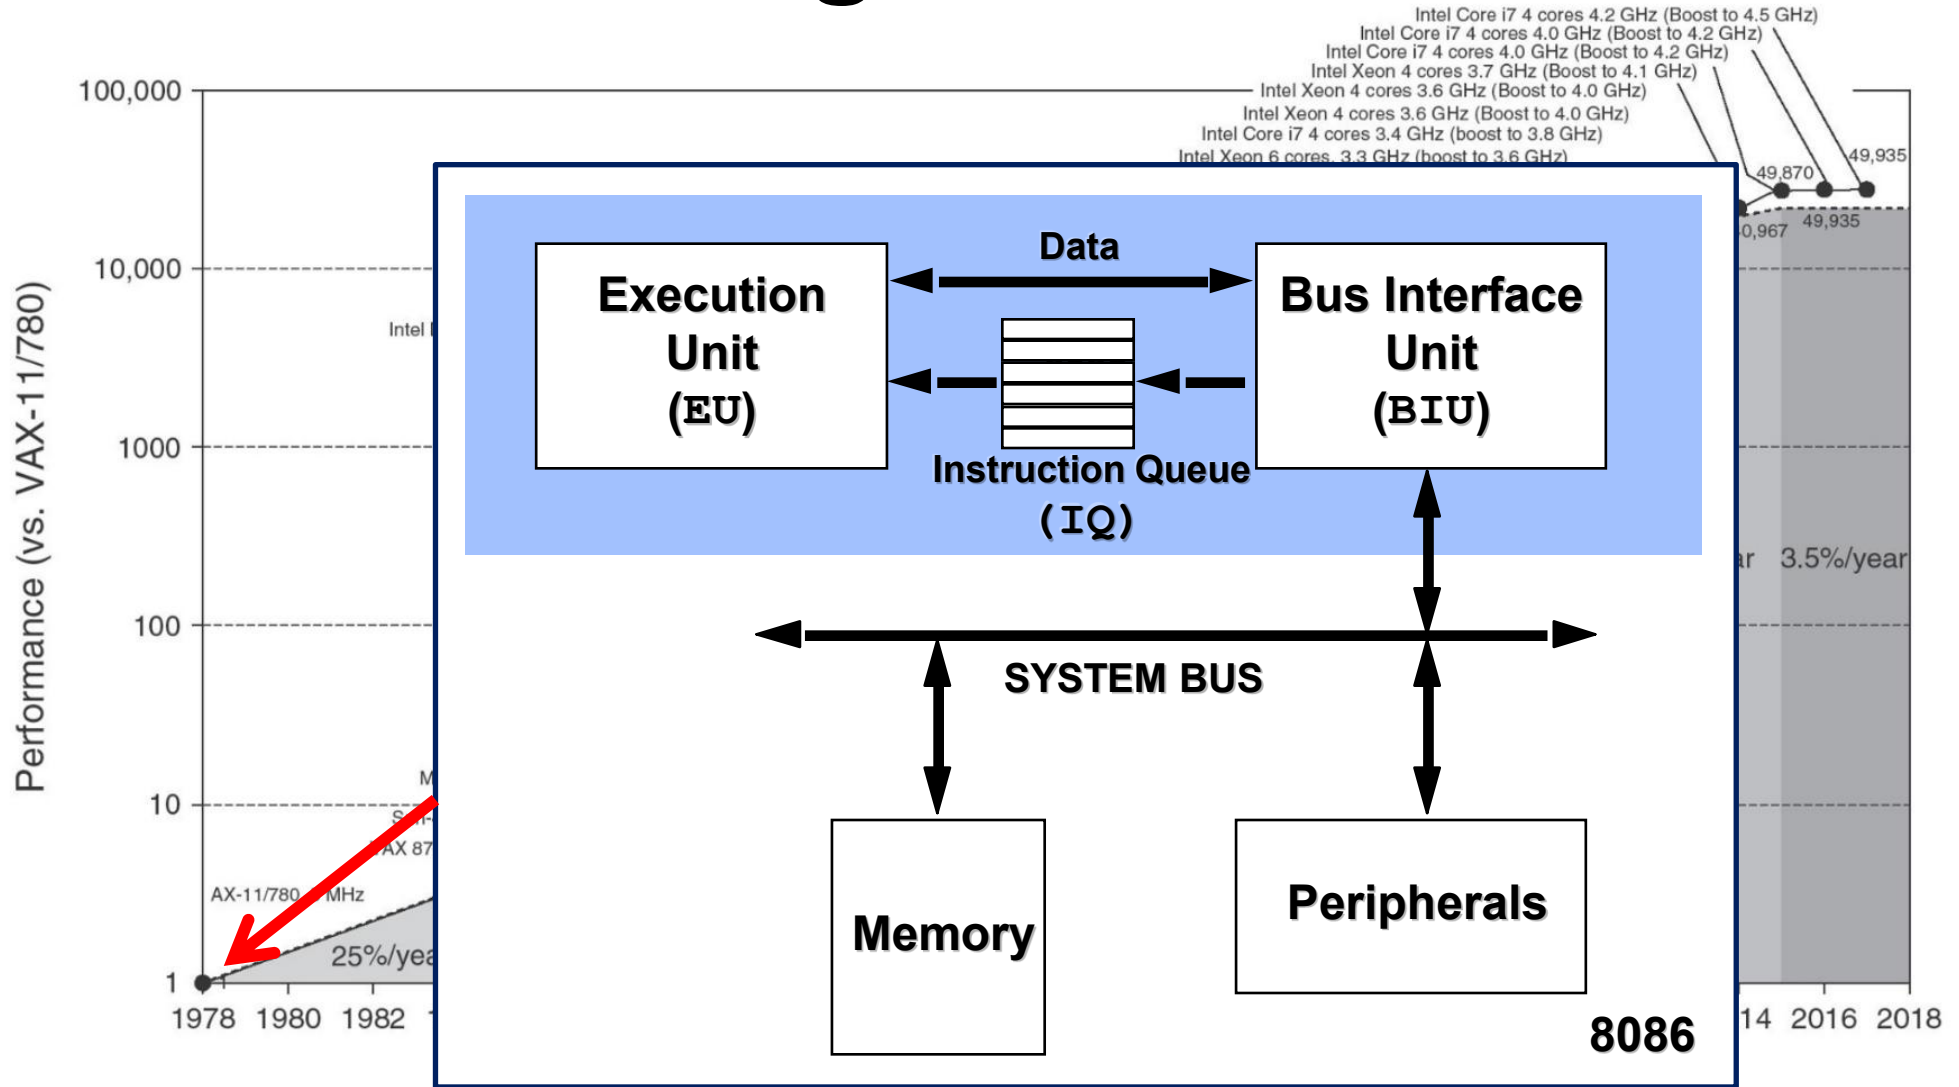
Microprocessor performance growth



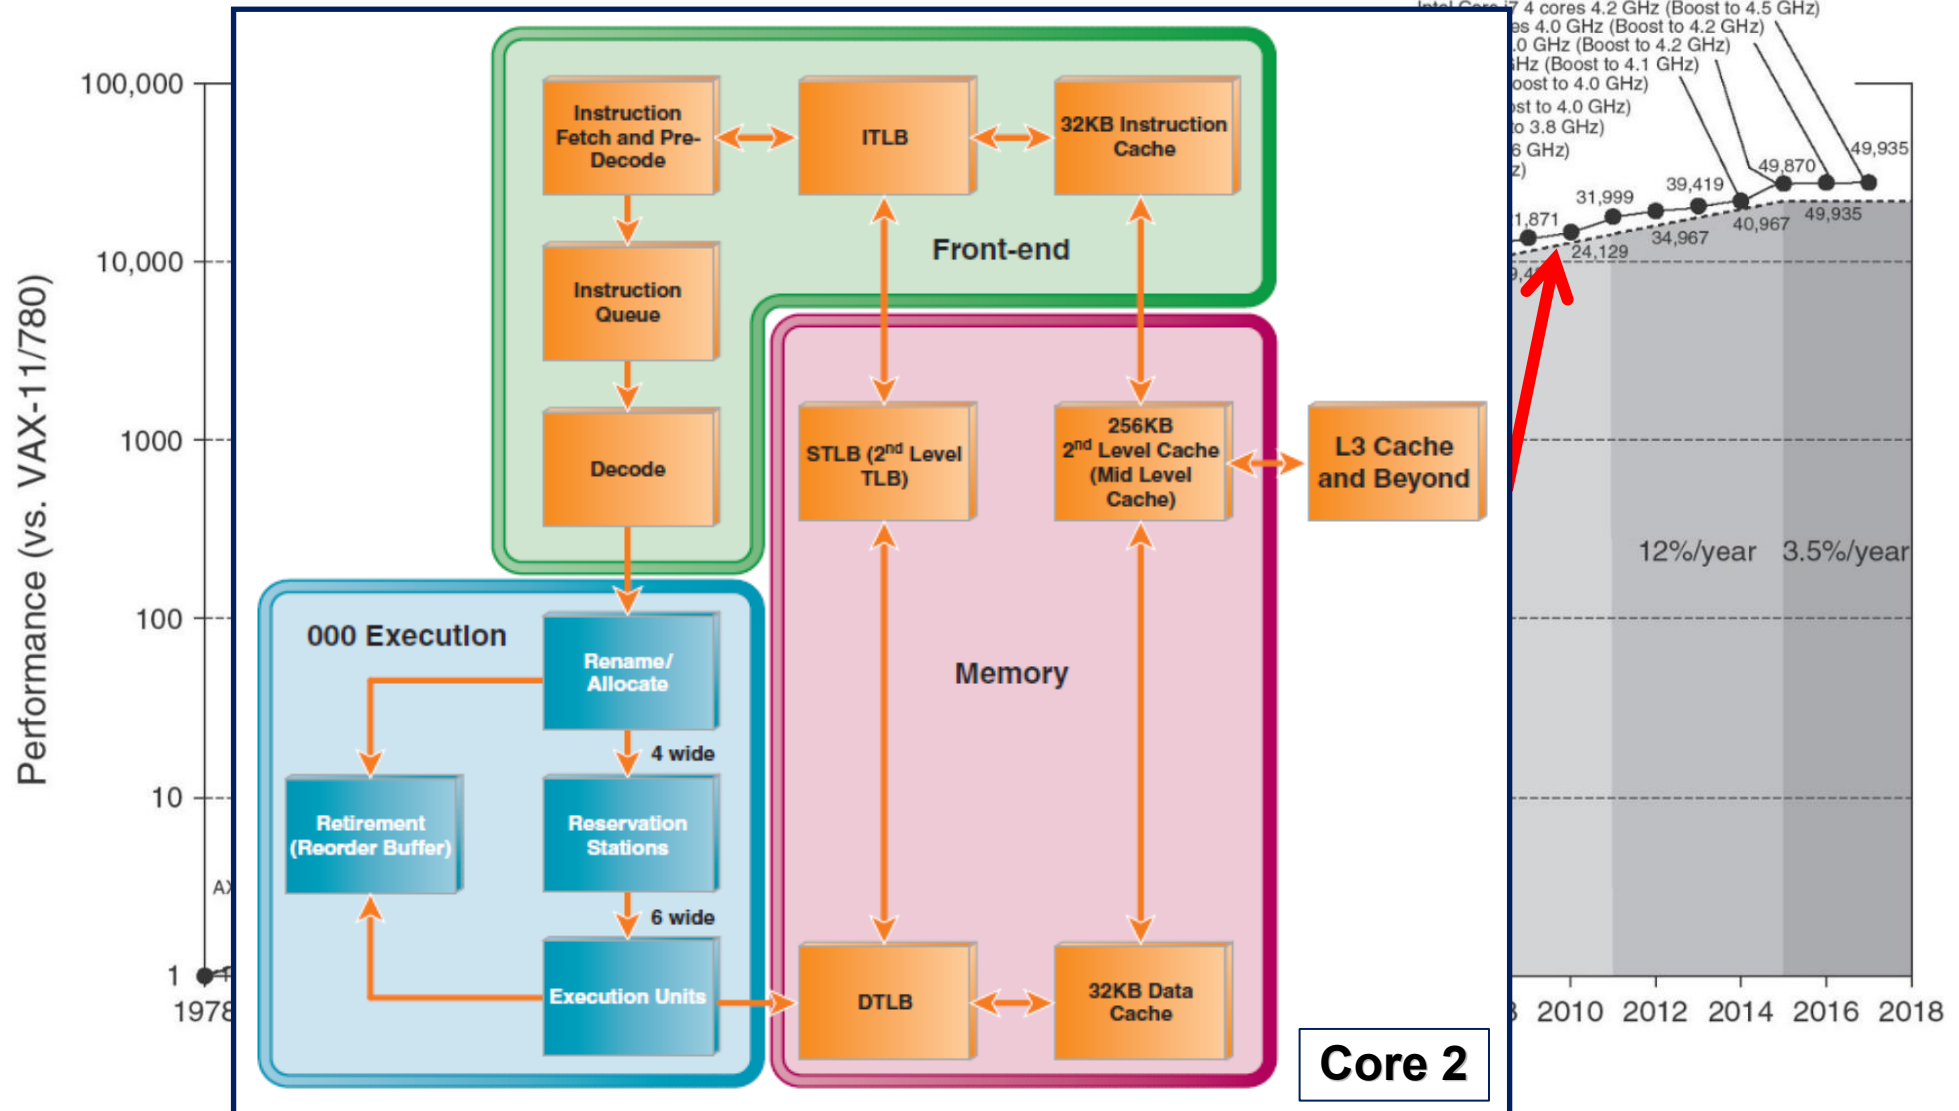
Microprocessor performance growth



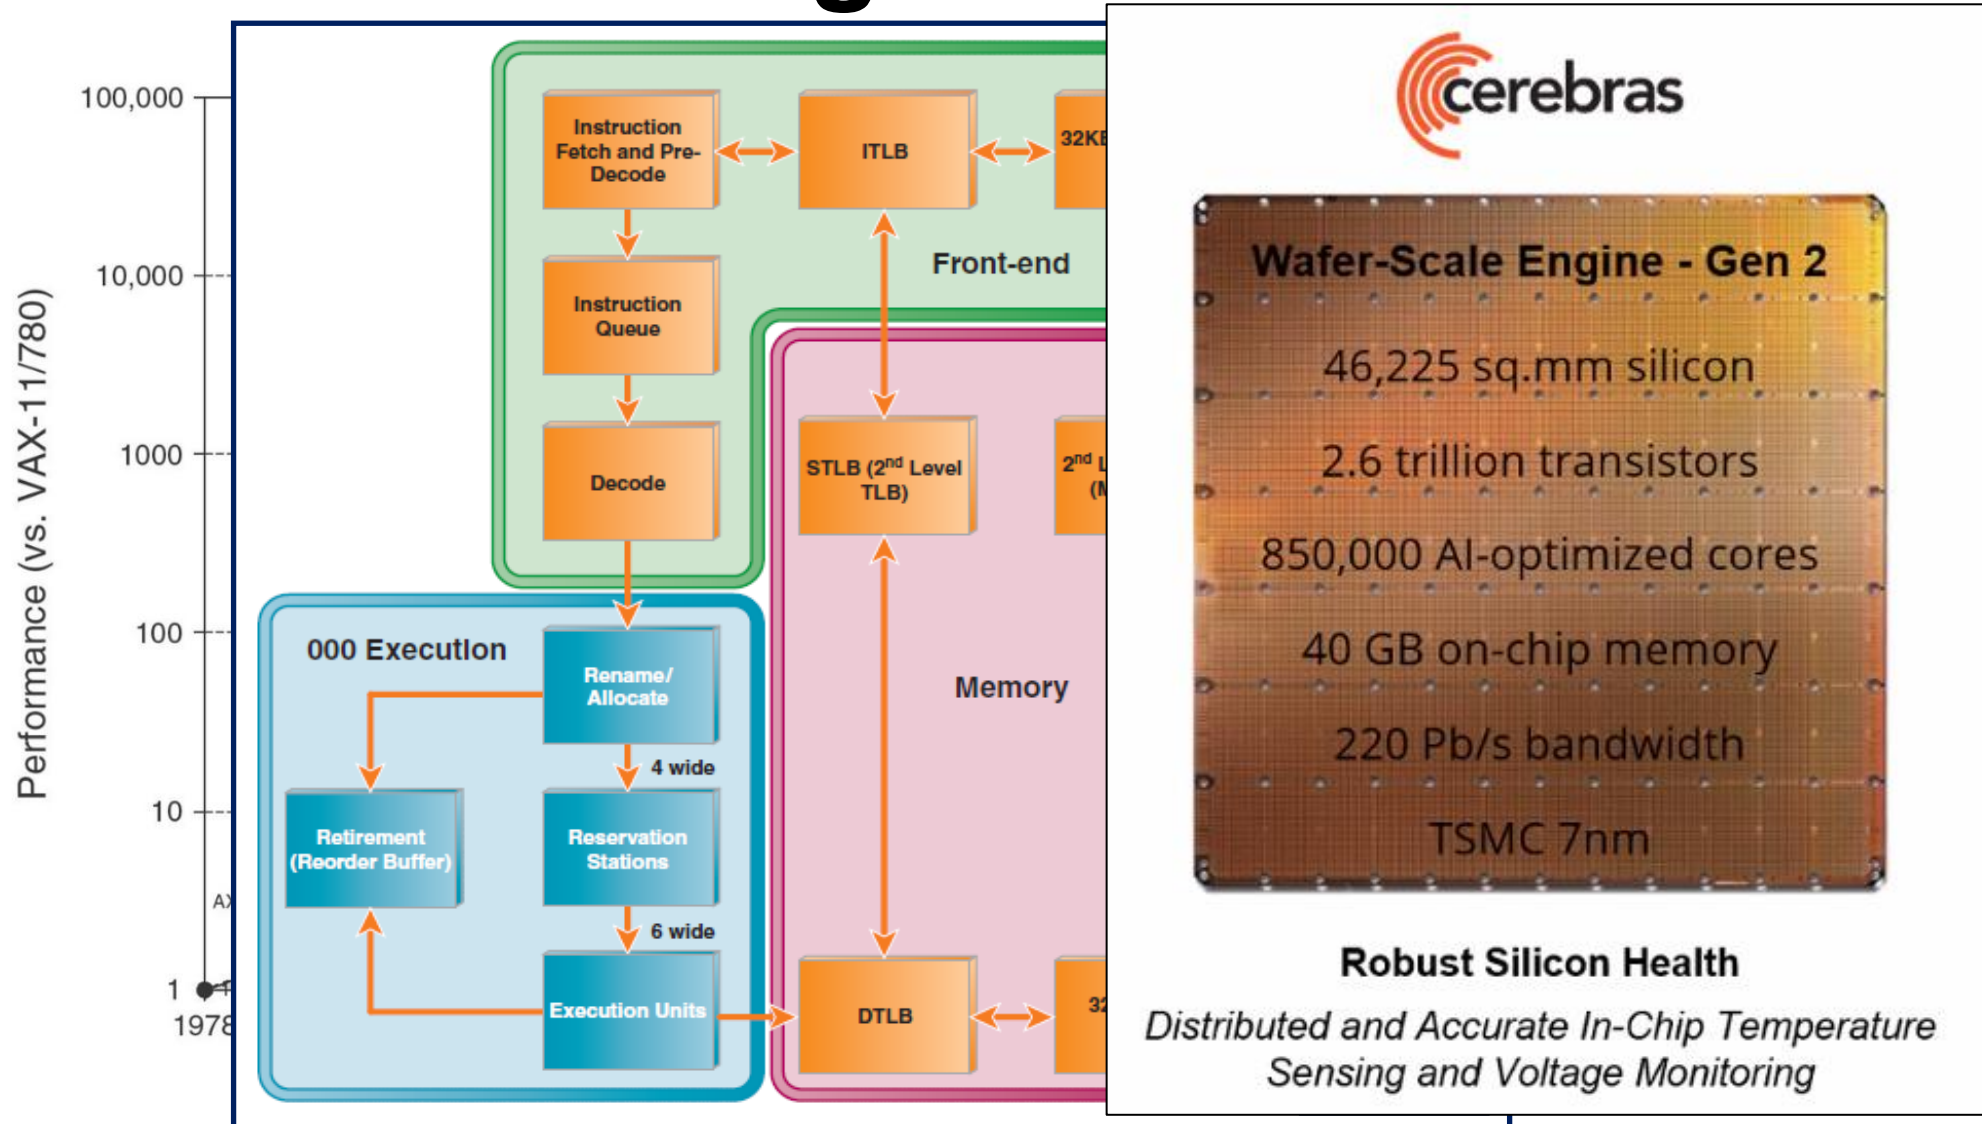
Microprocessor performance growth



Microprocessor performance growth



Microprocessor performance growth



Trend

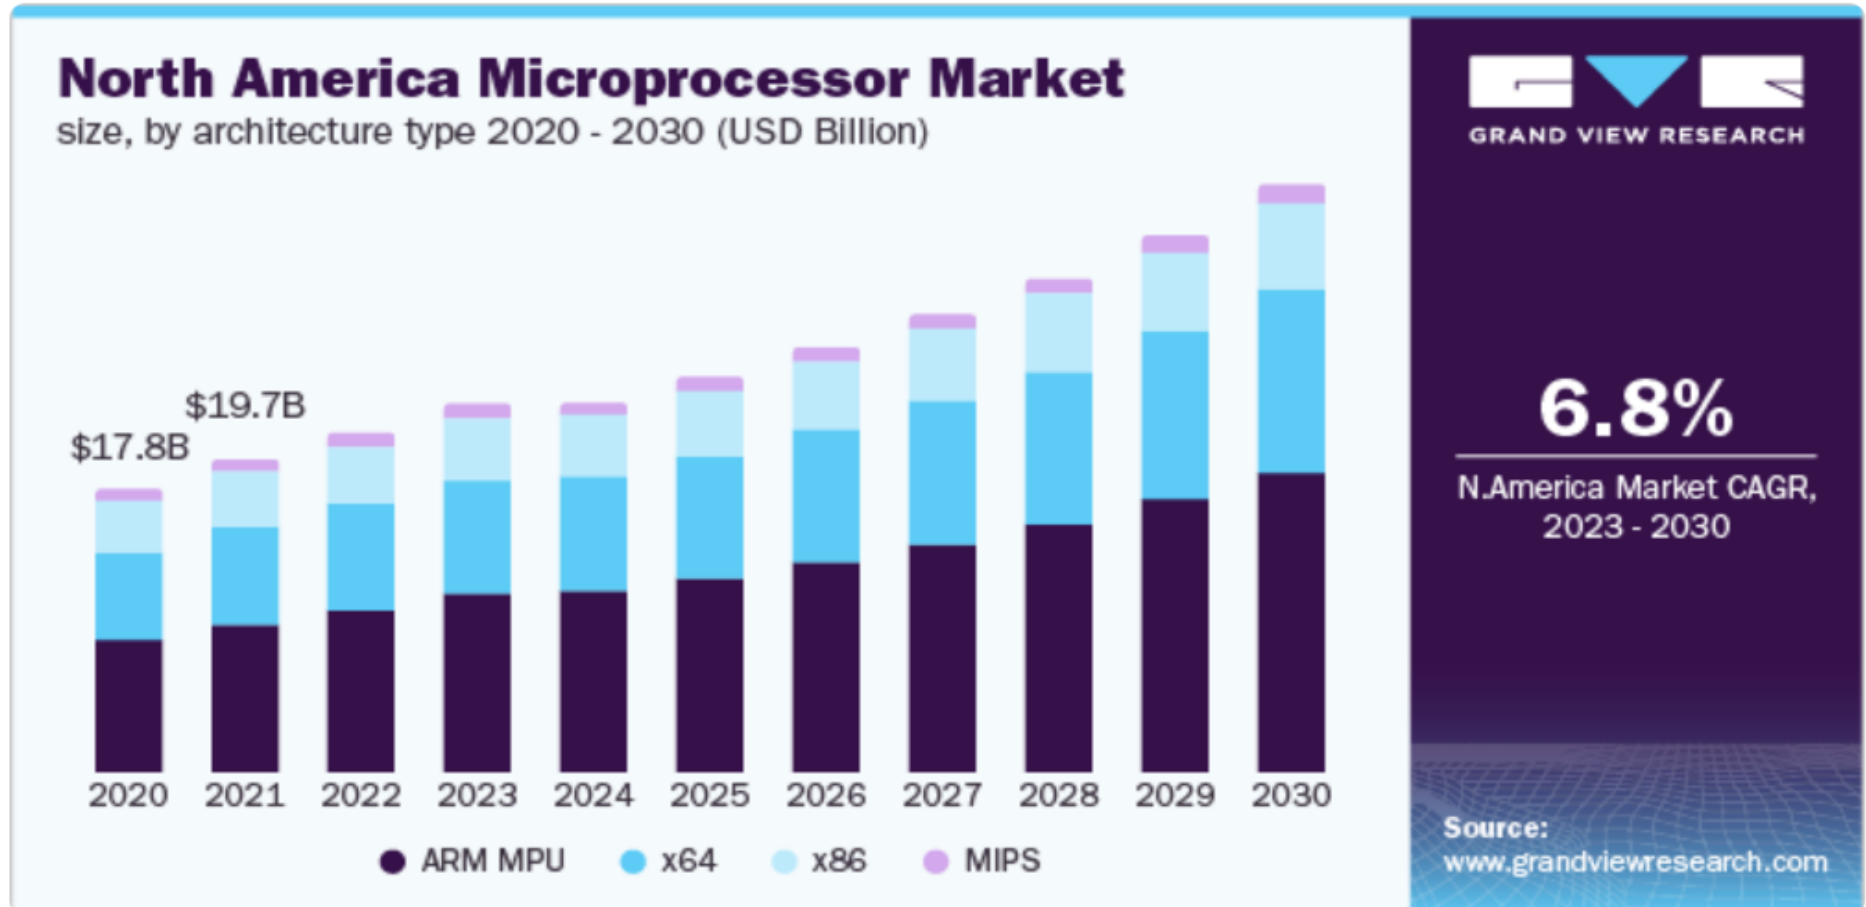
The major players in the microprocessor market (Intel, AMD, IBM, ARM) are not investing anymore in the development of faster processors, but rather on multiprocessor single chip systems (i.e., multicore devices).

The computer market


It is currently split in 5 different areas:

- **Personal Mobile Device (PMD)**
- **Desktop computing**
- **Servers**
- **Clusters / Warehouse Scale Computers (WSC)**
- **Embedded computers.**

Microprocessor market distribution



Microprocessor market distribution



Personal Mobile Device (PMD)

**This area includes smart phones, and tablet computers
Emphasis on energy efficiency and real-time applications.**

System price: \$100 - \$1,000

Microprocessor price: \$10 - \$100

Desktop computing

This area covers from PCs to workstations.

The main target of this area is to optimize the *price-performance* ratio.

System price: \$300 - \$2,500

Microprocessor price: \$50 - \$500

Servers

These systems provide larger-scale and more reliable computing services.

The main parameters of products in this area are:

- **Availability**
- **Scalability**
- **Throughput**

System price: \$5,000 - \$10,000,000

Microprocessor price: \$200 - \$2,000

Clusters / Warehouse-Scale Computers (WAS)

“Software as a Service (SaaS)”

“Platform as a Service (PaaS)”

Emphasis on availability, price-performance, and power consumption


Supercomputers, emphasis: floating-point performance and fast internal networks

System price: \$100,000 - \$200,000,000


Microprocessor price: \$50 - \$250

Supercomputers – HPC

The main difference between an HPC (High Performance Computers) and a conventional computer is its organization, interconnectivity, scale of electronic and software components.



HPC Performance



<https://www.top500.org/>

Embedded computers

This area is the fastest growing portion of the computer market.

It covers all special-purpose computer-based applications (from microwaves to coffee machines, from automotive to videogames)

Adopted microprocessors vary from cheap low-end 8-bit processors to very efficient (and expensive) high-end processors, but are not able to run third-party software.

System price: \$10 - \$100,000

Microprocessor price: \$0.01 - \$100

Embedded computers

Special requirements often existing in this area are

- **Real-time performance requirements**
- **Memory minimization**
- **Power consumption minimization**
- **Reliability constraints.**

Embedded computers

Embedded problems are often solved resorting to one of the following solutions:

- **Standard processor + custom logic + custom SW**
- **Standard processor + custom SW**
- **Standard DSP + custom SW.**

Programmable devices (FPGAs) are playing a growing role.

Classes of Parallelism

- **Data-level Parallelism (DLP):**

Many data items that can be operated on at the same time

- **Task-level Parallelism (TaskLP):**

Different tasks of a work can operate independently.

Parallel Architectures

- **Instruction-level Parallelism (ILP):**
Modestly exploits Data-level Parallelism
- **Vector Architectures and Graphic processor unit (GPUs):**
Exploits Data-level Parallelism
- **Thread-level Parallelism (TLP):**
Exploits Data-level Parallelism and Task-level Parallelism
- **Request-level Parallelism (RLP):**
Exploits parallelism among decoupled tasks.

Designing a computer

It means

- **determining which attributes are important for the new machine**
- **designing a machine which**
 - **maximizes performance and**
 - **matches cost and power constraints.**

Computer architecture

In the last decades, computer design took advantage of

- **Architectural innovation**
- **Technology improvements**
- **Software development.**

It was estimated that the difference between the highest-performance microprocessors available in 2001 and what would have been obtained by relying solely on technology is more than a factor of 15.

Computer architecture

It includes three aspects of computer design:

- **Instruction set architecture**
- **Organization**
- **Hardware.**


Computer architect must design a computer meeting:

- **Functional requirements**
- **Price**
- **Power**
- **Performance**
- **Dependability.**

Moore's Law

The number of devices (i.e., transistors) that can be integrated into a single chip doubles every 18/24 months.

Intel processors complexity growth



IC manufacturing cost

When evaluating it, it is important not to forget the impact of *yield*, i.e., the percentage of products that pass the test phase.

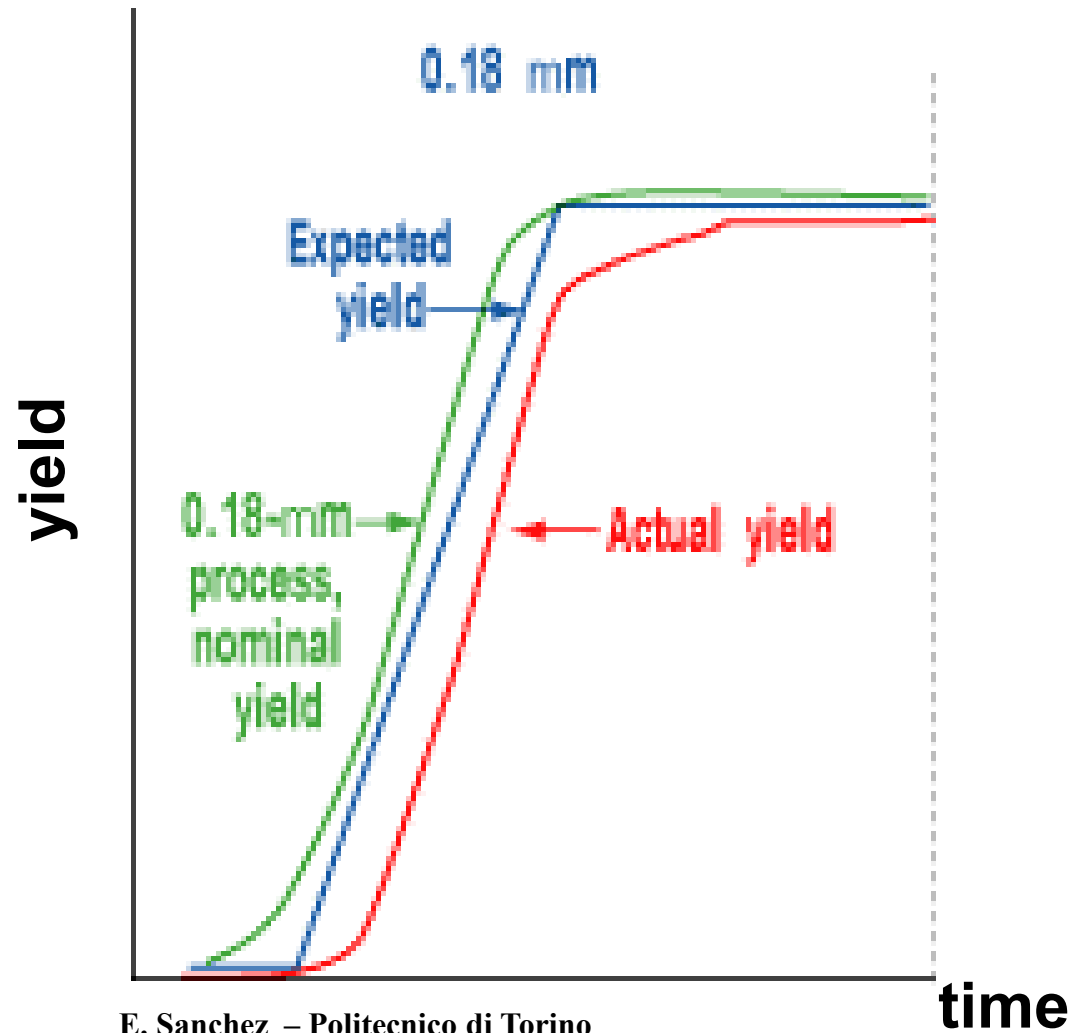
The production process for every product undergoes an evolution which normally leads to an improvement in yield (also known as *learning curve*).

When yield increases, the cost decreases.

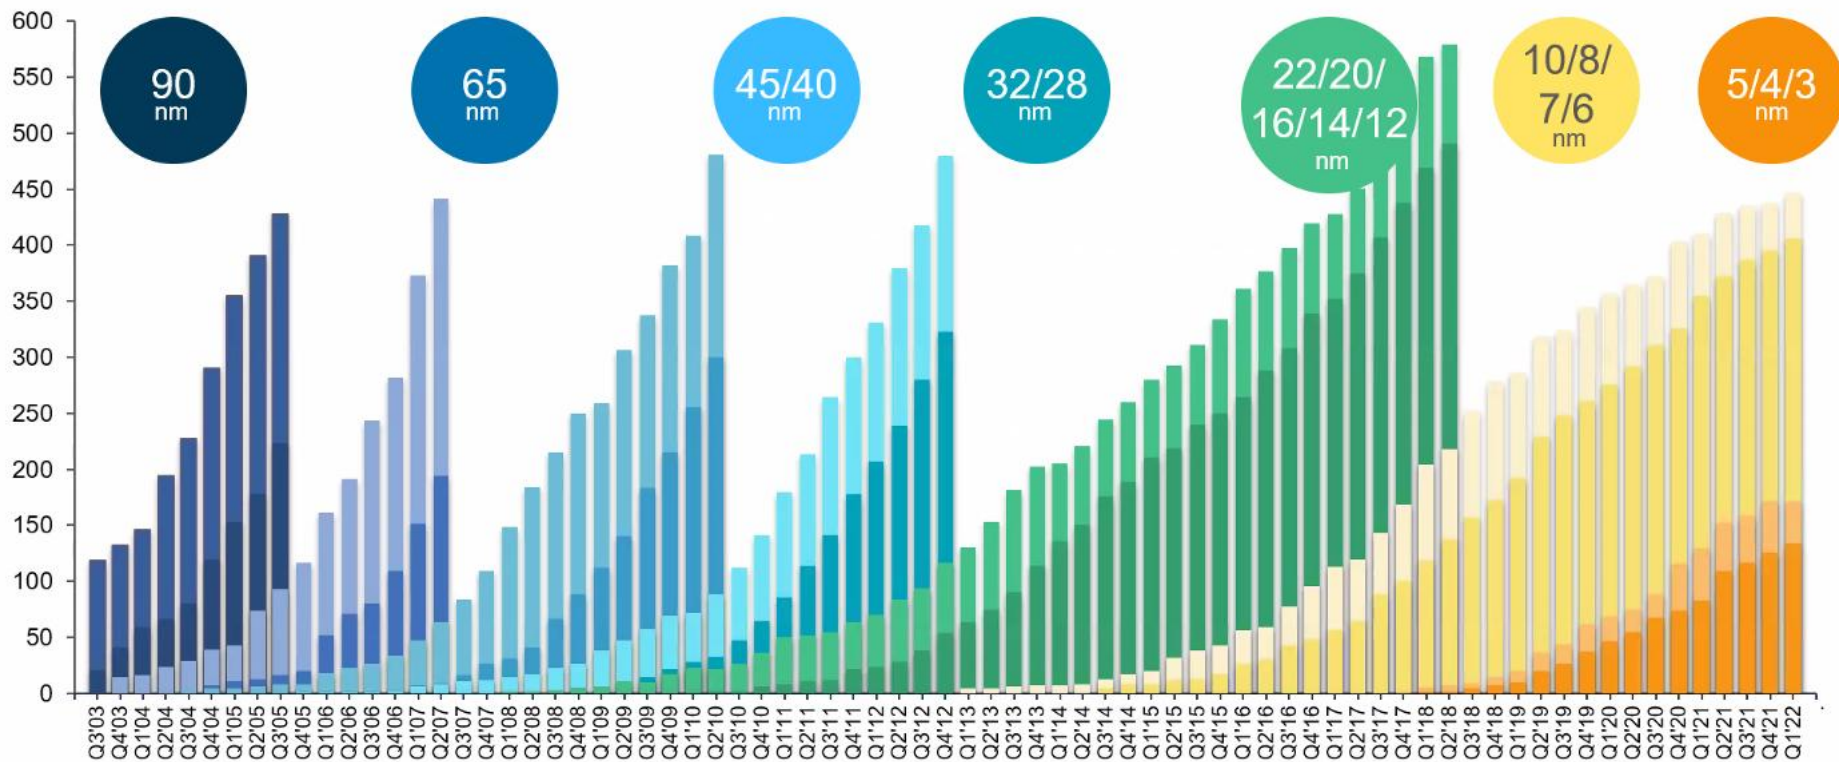


More than 50% of manufacturing cost is due to validation and testing procedures!!

Yield behavior



Industry adoption of new transistor technologies



Power consumption

Continuous increase in system complexity and device integration often lead to problems with power consumption.

Power consumption may be critical under two aspects:

- **Power (static and dynamic)**
- **Energy (mainly for portable devices).**

Power

Until now it has been dominated by dynamic power, i.e., that consumed by each transistor when switching between different states.

The dynamic power for each transistor can be evaluated by the formula

$$\text{Power}_{\text{dynamic}} = \frac{1}{2} \times \text{capacitive load} \times \text{voltage}^2 \times \text{frequency}$$

$$\text{Power}_{\text{static}} = V \times I \quad (25 - 50\% \text{ of total power consumption})$$

For this reason, voltage continuously dropped in the last years.

Energy

It is given by

$$\text{Energy}_{\text{dynamic}} = \text{capacitive load} \times \text{voltage}^2$$

It is mainly of interest for mobile devices.

Dependability

Dependability is the quality of the system to deliver a correct service.

Dependability of computer systems is traditionally very high, but it can be lowered by

- **Bugs in the design of the hardware**
- **Bugs in the software**
- **Defects in the hardware (introduced by the manufacturing process)**
- **Faults happening during the product operation.**

Are DNN applications safe?

FINANCIAL TIMES

Worker at Volkswagen plant killed in robot accident

Fatality touches on concerns about spread of automation



Chris Bryant in Frankfurt and Richard Waters in San Francisco JULY 2 2015

A technician was killed by a robot at a [Volkswagen](#) plant near Kassel, Germany on Wednesday, in a rare accident that touches on concerns about the spread of automation and its impact on jobs.


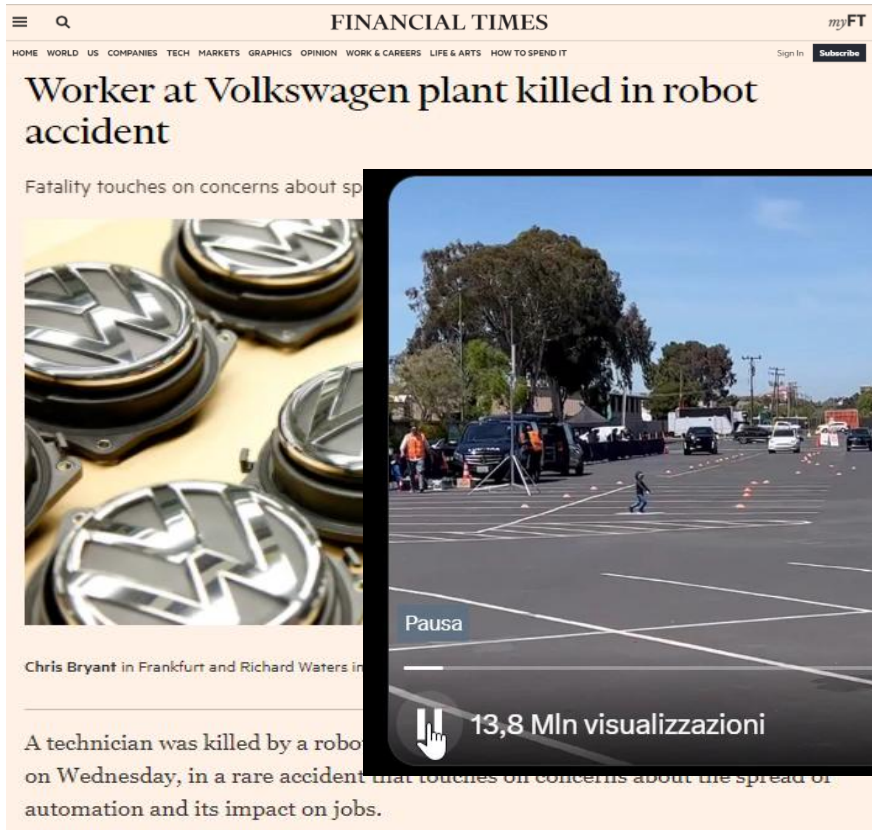
The New York Times

Tesla Autopilot System Found Probably at Fault in 2018 Crash

The National Transportation Safety Board called for improvements in the electric-car company's driver-assistance feature and cited failures by other agencies.



Are DNN applications safe?



Safety-critical applications

In the past, possible misbehaviors of computer systems were critical in some areas, such as

- **Space**
- **Avionics**
- **Nuclear plants control.**

In more recent years, computer-based systems expanded to other safety-critical areas, such as

- **Rail-road traffic control**
- **Automotive**
- **Biomedical**
- **Telecommunications.**

Dependability importance

It increased very much, because in several areas it is crucial to guarantee that the system matches the dependability constraints, e.g., in terms of probability of behaving as expected for long periods.

Dependability evaluation

Dependability is often measured using

- Mean Time To Failure (MTTF) or Failures In Time (FIT), which is its reciprocal. 1 FIT = 1 failure in one billion hours
- Mean Time Between Failures (MTBF)
- Mean Time To Repair (MTTR)


The three measures are related by the formula

$$\text{MTBF} = \text{MTTF} + \text{MTTR}$$

Finally, availability is the probability that a system works correctly in a generic time instant.

The Bathtub Curve

Hypothetical Failure Rate versus Time



Computer performance

What is performance?

It can be measured by counting the number of operations executed or the time needed to execute them.

User point of view:

performance = *response time* (time between start and completion of an operation)

System manager point of view:

performance = *throughput* (total amount of work done in the time unit).

Time

Which time has to be considered for performance computation?

- **Elapsed time**
- **CPU time**
 - **user CPU time**
 - **system CPU time.**

All these measures can be of interest. UNIX provides all of them through the `time` command

90.7s 12.9s 2:39 65%


Time

Which time has to be considered for performance computation?

- Elapsed time
- CPU time
 - user CPU time
 - system CPU time.

All these measures can be of interest. UNIX provides all of them through the `time` command

90.7s 12.9s 2:39 65%



Performance evaluation

It is often performed by letting the computer to execute applications and observing its behavior.

Unfortunately, the choice of the applications severely affects the performance.

In the ideal case, one should use as *workload* the mix of applications the user will run.

However, they are normally unknown, and largely variable from one user to another.

Therefore, some *benchmarks* are selected to mimic real cases.

Program benchmarks

Possible benchmarks:

- ***Real programs*** (e.g., C compilers, text processors, special-purpose tools), possibly modified
- ***Kernels*** (e.g., Livermore Loops, Linpack)
- ***Toy benchmarks*** (e.g., Quicksort, Sieve of Eratosthenes)
- ***Synthetic benchmarks*** (e.g., Whetstone, Dhrystone).

Benchmark suites

They contain a number of different programs, so that the weakness of any component is lessened by the presence of the others.

Benchmark sets are normally composed of:

- **kernels**
- **program fragments**
- **applications.**

SPEC evolution

Standard
Performance
Evaluation
Corporation

	Benchmark name by SPEC generation					
	SPEC2017	SPEC2006	SPEC2000	SPEC95	SPEC92	SPEC89
GNU C compiler	←					gcc
Perl interpreter	←					espresso
Route planning	←		mcf			li
General data compression	XZ		bzip2		compress	eqntott
Discrete Event simulation - computer network	←	omnetpp	vortex	go	sc	
XML to HTML conversion via XSLT	←	xalancbmk	gzip	jpeg		
Video compression	X264	h264ref	eon	m88ksim		
Artificial Intelligence: alpha-beta tree search (Chess)	deepsjeng	sjeng	twolf			
Artificial Intelligence: Monte Carlo tree search (Go)	leela	gobmk	vortex			
Artificial Intelligence: recursive solution generator (Sudoku)	exchange2	astar	vpr			
		hmmer	crafty			
		libquantum	parser			
Explosion modeling	←	bwaves				fpppp
Physics: relativity	←	cactuBSSN				tomcatv
Molecular dynamics	←	namd				doduc
Ray tracing	←	povray				nasa7
Fluid dynamics	←	lbm				spice
Weather forecasting	←	wrf			swim	matrix300
Biomedical imaging: optical tomography with finite elements	parest	gamess		apsi	hydro2d	
3D rendering and animation	blender			mgrid	su2cor	
Atmosphere modeling	cam4	milc	wupwise	applu	wave5	
Image manipulation	imagick	zeusmp	apply	turb3d		
Molecular dynamics	nab	gromacs	galgel			
Computational Electromagnetics	fotonik3d	leslie3d	mesa			
Regional ocean modeling	roms	dealll	art			
		soplex	equake			
		calculix	facerec			
		GemsFDTD	ammp			
		tonto	lucas			
		sphinx3	fma3d			
			sixtrack			

MiBench Benchmarks

Auto./Industrial	Consumer	Office	Network	Security	Telecomm.
basimath	jpeg	ghostscript	dijkstra	blowfish enc.	CRC32
bitcount	lame	ispell	patricia	blowfish dec.	FFT
qsort	mad	rsynth	(CRC32)	pgp sign	IFFT
susan (edges)	tiff2bw	sphinx	(sha)	pgp verify	ADPCM enc.
susan (corners)	tiff2rgba	stringsearch	(blowfish)	rijndael enc.	ADPCM dec.
susan (smoothing)	tiffdither			rijndael dec.	GSM enc.
	tiffmedian			sha	GSM dec.
	typeset				

Reproducibility

Information about execution times on benchmarks should allow reproducibility.

This means reporting detailed information about

- **hardware (system configuration)**
- **software (OS, compiler, program)**
- **program input.**

Comparing and summarizing performance

Problem 1

I know the performance of one machine on a set of programs: which is its global performance?

Problem 2

I know the performance of two machines on the same set of programs: which is their relative performance?

A number of metrics have been proposed.

Index i runs
over all
benchmarks
in the set

Total execution time

$$\sum_{i=1}^n \text{Time}_i$$

Normalized execution time

A reference machine is adopted (e.g., VAX-11/780) and execution times are normalized with respect to it.

Arithmetic Mean

Arithmetic mean:

$$\frac{1}{n} \sum_{i=1}^n \text{Time}_i$$

Weighted mean

Weighted arithmetic mean:

$$\sum_{i=1}^n \text{Weight}_i \times \text{Time}_i$$

Suggested solution

To measure a real workload and weight the programs according to their frequency of execution.

Program inputs should be carefully specified.

Example of weighted arithmetic mean execution time

	Machines			Weightings		
	A	B	C	W(1)	W(2)	W(3)
Program P1 (secs)	1.00	10.00	20.00	0.50	0.909	0.999
Program P2 (secs)	1000.00	100.00	20.00	0.50	0.091	0.001
Arithmetic mean: W(1)	500.50	55.00	20.00			
Arithmetic mean: W(2)	91.91	18.19	20.00			
Arithmetic mean: W(3)	2.00	10.09	20.00			

Guidelines and Principles for Computer Design

- Amdahl's law
- CPU performance equation.

Amdahl's Law: preliminaries

$$\text{speedup} = \frac{\text{performance with enhancement}}{\text{performance without enhancement}}$$

The speedup resulting from an enhancement depends on two factors:

- **$\text{fraction}_{\text{enhanced}}$** : the fraction of the computation time that takes advantage of the enhancement
- **$\text{speedup}_{\text{enhanced}}$** : the size of the enhancement on the parts it affects.

Amdahl's Law

execution time_{new} =

$$\text{execution time}_{\text{old}} \times \left((1 - \text{fraction}_{\text{enhanced}}) + \frac{\text{fraction}_{\text{enhanced}}}{\text{speedup}_{\text{enhanced}}} \right)$$

$$\text{speedup}_{\text{overall}} = \frac{\text{execution time}_{\text{old}}}{\text{execution time}_{\text{new}}} =$$

1

$$(1 - \text{fraction}_{\text{enhanced}}) + \frac{\text{fraction}_{\text{enhanced}}}{\text{speedup}_{\text{enhanced}}}$$

Example

An enhancement makes one machine 10 times faster for 40% of the programs the machine runs. Which is the overall speedup?

$$\text{fraction}_{\text{enhanced}} = 0.4$$

$$\text{speedup}_{\text{enhanced}} = 10$$

$$\text{speedup}_{\text{overall}} = \frac{1}{(1 - 0.4) + \frac{0.4}{10}} = 1.56$$

Choosing between two solutions: an example

Two solutions are available for increasing the floating-point performance of one machine.

Solution 1

Increasing by 10 the performance of square root operations (responsible for 20% of the execution time) by adding specialized hardware.

Solution 2

Increasing by 2 the performance of all the floating-point operations (responsible for 50% of the execution time).

Which solution makes the machine faster?

Amdahl's Law application

Solution 1

$$\text{speedup}_1 = \frac{1}{(1 - 0.2) + \frac{0.2}{10}} = 1.22$$

Solution 2

$$\text{speedup}_2 = \frac{1}{(1 - 0.5) + \frac{0.5}{2}} = 1.33$$

Measuring the time required to execute a program

Possible approaches:

- by observing the real system
- by simulation
- by applying the CPU performance equation.

The CPU Performance Equation

$$\text{CPU time} = \left(\sum_{i=1}^n \text{CPI}_i \times \text{IC}_i \right) \times \text{Clock cycle time}$$

where

- **CPI_i** is the number of clock cycles required by instruction i
- **IC_i** is the number of times instruction i is executed in the program
- ***clock cycle time*** is the inverse of clock frequency.

Timing Equation

Depends on the hardware organization and instruction set architecture

clock cycle time

where

- CPI_i is the number of clock cycles required by instruction i
- IC_i is the number of times instruction i is executed in the program
- *clock cycle time* is the inverse of clock frequency.

The CPU Performance Equation

CPU time

n

time

Depends on the instruction set
architecture and compiler
technology

where

- CPI_i is the number of instructions executed by instruction i
- IC_i is the number of times instruction i is executed in the program
- *clock cycle time* is the inverse of clock frequency.

The CPU Performance Equation

$$\text{CPU time} = \left(\sum_{i=1}^n \text{CPI}_i \times \text{IC}_i \right) \times \text{Clock cycle time}$$

where

- CPI_i is the number of clock cycles per instruction by the processor. It depends on the hardware technology and organization.
- IC_i is the number of instructions executed in the program.
- *clock cycle time* is the inverse of clock frequency.

Depends on the hardware technology and organization

The CPU Performance Equation limitations

In pipelined processors, CPI_i may vary for a given instruction, depending on different parameters

- Instructions executed before and after**
- Memory system behavior (e.g., cache miss or hit)**

Therefore, evaluating the execution time analytically becomes much harder.