

Branch Prediction Techniques

E. Sanchez

**Politecnico di Torino
Dipartimento di Automatica e Informatica**

INTRODUCTION

Branches can potentially impact in a very serious way the pipeline performance.

It is possible to reduce the performance losses by predicting how branches will behave.

Branch prediction schemes

They aim at correctly forecasting branches, thus reducing the chance that control dependences cause stalls.

They can be categorized in two groups:

- *static techniques*: they are handled by the compiler resorting to a preliminary analysis of the code
- *dynamic techniques*: they are implemented by the hardware based on the behavior of the code.

Static branch prediction

It can be useful when combined with other static techniques, such as

- Enabling delayed branches
- Rescheduling to avoid data hazards.

Static branch prediction

The compiler may predict branch behavior in different alternative ways:

- Always predicting branches as taken
- Predicting branch behavior depending on branch direction
- Predicting on the basis of profile information coming from earlier runs.

Predicting branches as taken

On the SPEC92 programs this gives:

- 34% average misprediction rate
- highly variable rate (from 9% to 59%).

Other techniques may behave better in the average, but still with very high variations from program to program.

Predicting branch behavior depending on branch direction

This technique is based on the observation that

- Forward branches are more often untaken
- Backward branches are more often taken.

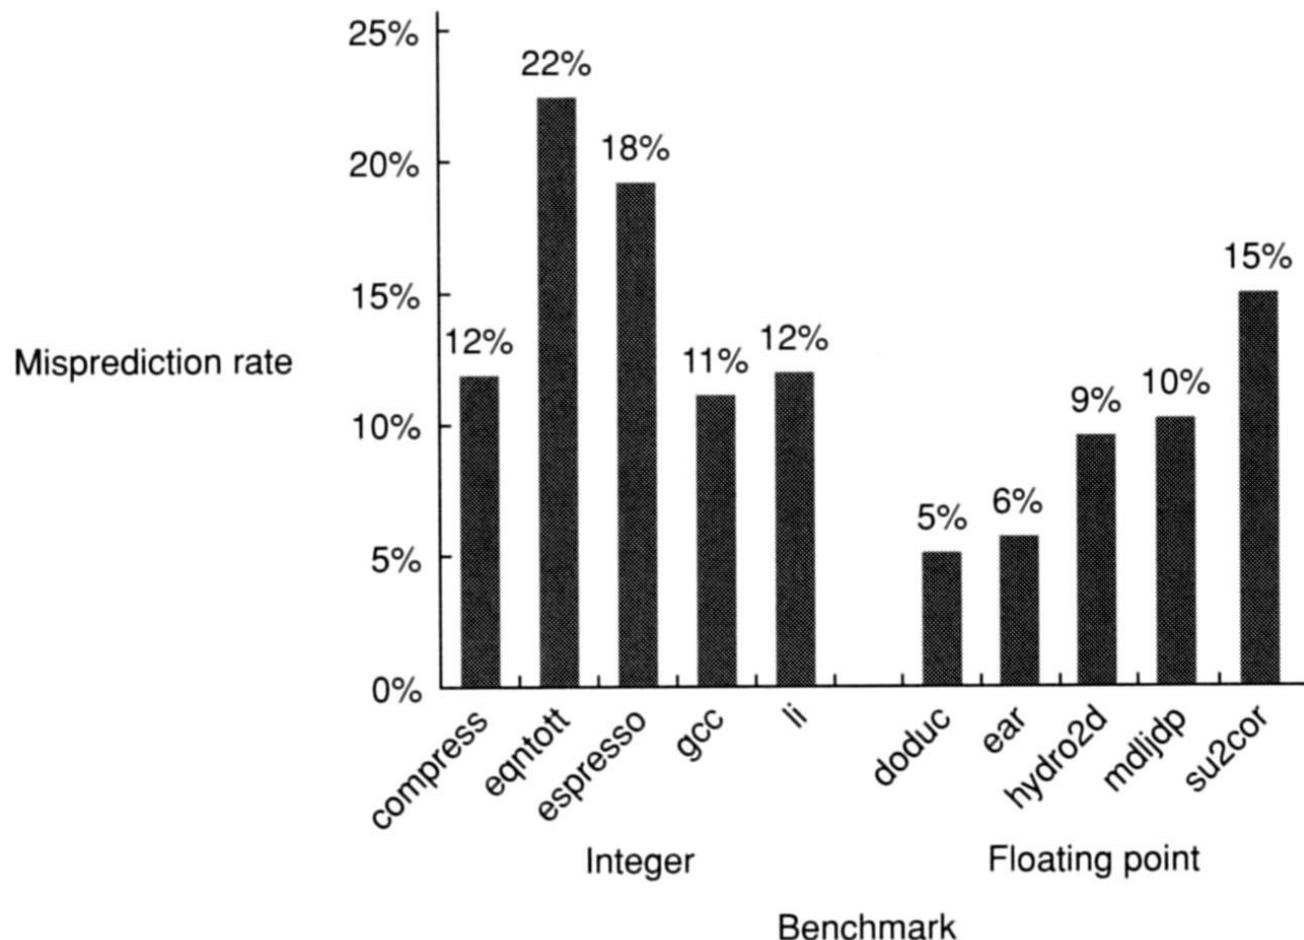
This behavior is mainly due to loop constructs.

Predicting branches based on profile information

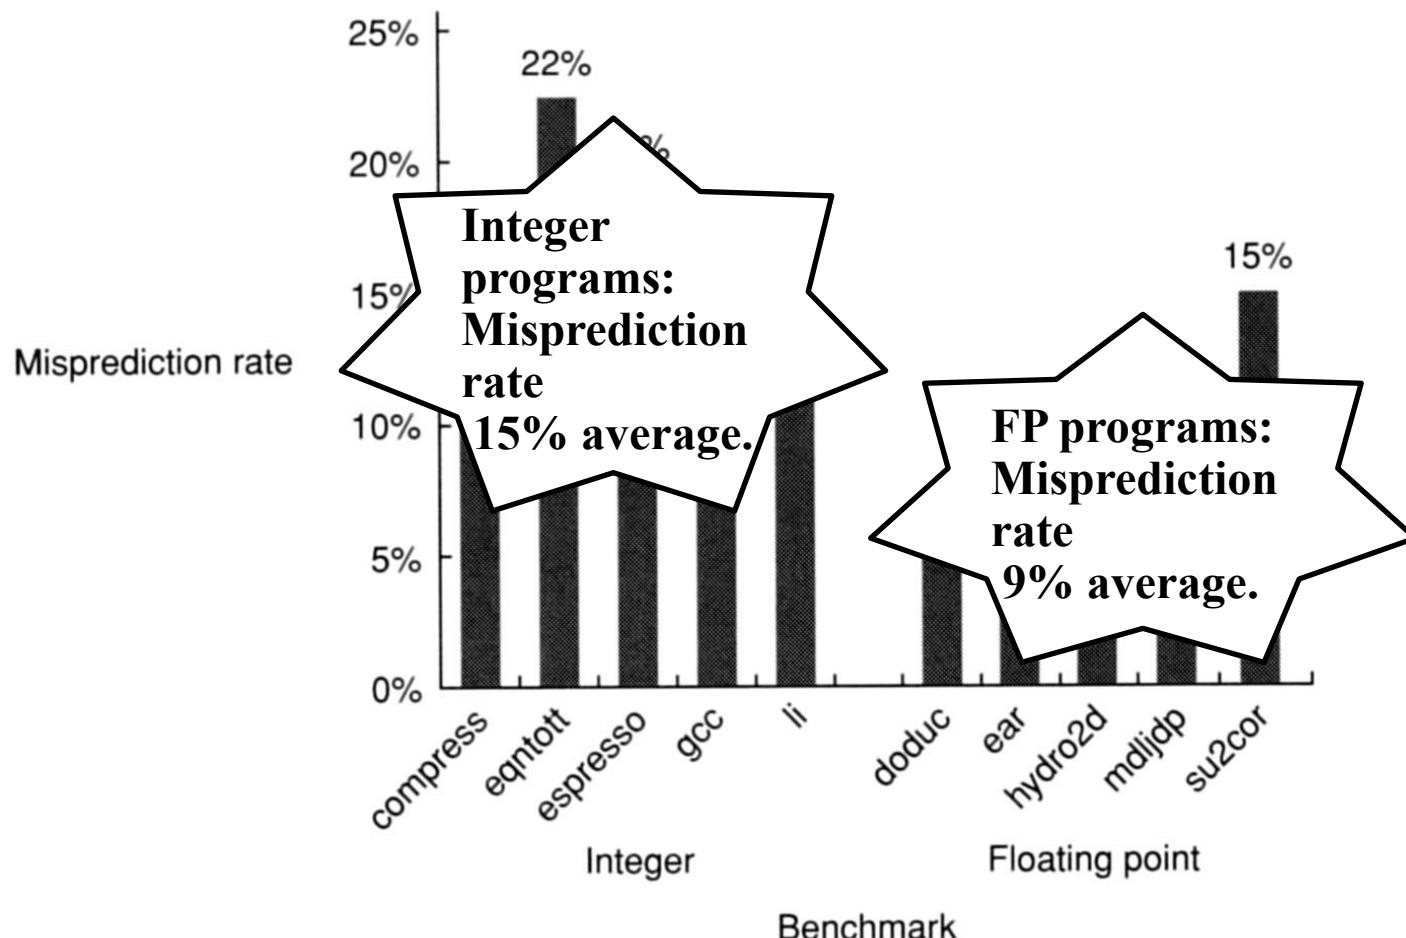
This technique is based on

- Identifying a typical sequence of input stimula for the program
- Running the program a limited number of times using these stimula
- Gathering statistics on the behavior of branches
- Using these statistics as predictions.

Profile-based prediction effectiveness



Profile-based prediction effectiveness



Dynamic branch prediction

Dynamic schemas are based on hardware, and use the branch instruction address to activate the different prediction mechanisms.

Dynamic branch prediction can be based on different techniques:

- **Branch history table**
 - **One- and two-bit prediction schemes**
- **Two-level prediction schemes (correlating predictors)**
- **Branch-target buffer**
- **Others.**

BRANCH HISTORY TABLE

It is the simplest method for dynamic branch prediction.

The *Branch History Table* (BHT) is a small memory:

- indexed by the lower portion of the address of the branch instruction
- containing for each entry one or more bits recording whether the branch has been taken or not the last time it has been executed.

Algorithm

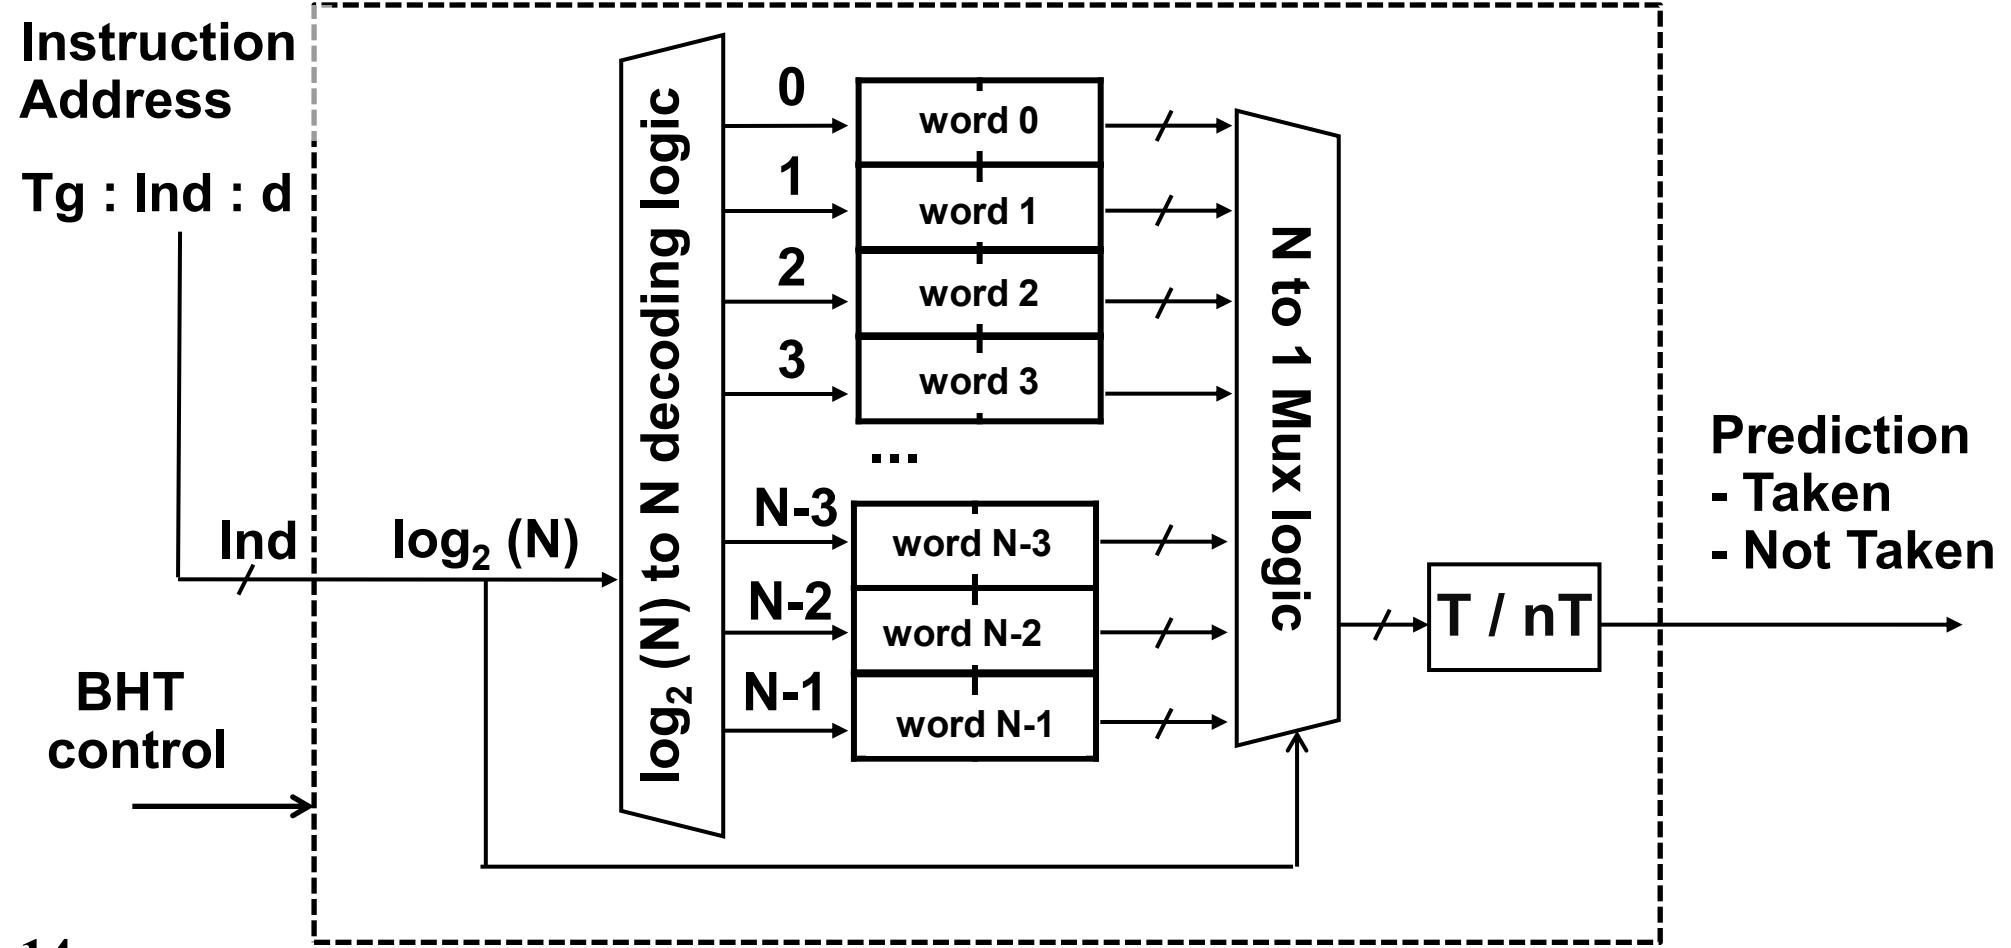
Each time a branch is decoded

- An access is made in the BHT using the lowest portion of its address as an index
- The prediction stored in the table is used, and the new PC is computed according to the prediction.

When the branch result is known, the BHT is possibly updated.

A BHT using only one bit at every table location is called a BIMODAL predictor.

BHT implementation



Effectiveness

The effectiveness of the method depends on:

- the chance that the BHT entry relates to the branch of interest
- the accuracy of the prediction.

In the MIPS processor, the evaluation of the branch condition is performed while branch instructions are identified. Therefore, the described technique does not give any advantage.

Example

Consider a loop branch which is taken nine times in a row, then not taken once. Assume that the entry for the branch is not shared with other branches.

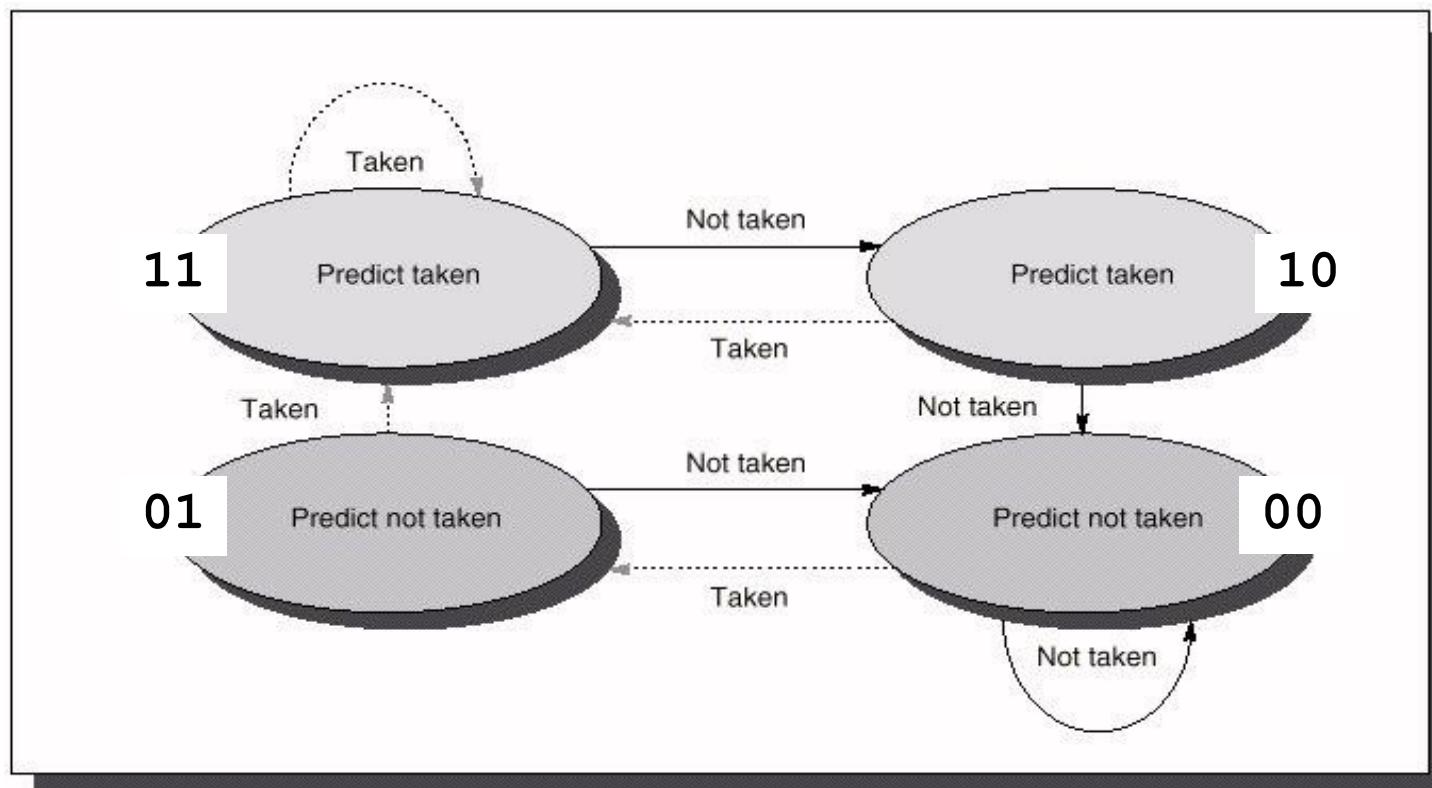
The steady-state prediction behavior using a 1-bit BHT will mispredict on the first and last loop iterations.

The prediction accuracy is thus 80%, lower than simply assuming that the branch is taken (90%).

Two-bit Prediction Schemes

They provide higher prediction capabilities.

For every branch, two bits are maintained, and the prediction is changed only after missing twice.



n-bit Prediction Scheme

It is the general case of the previous one.

It is based on an n-bit saturating counter associated to every branch instruction.

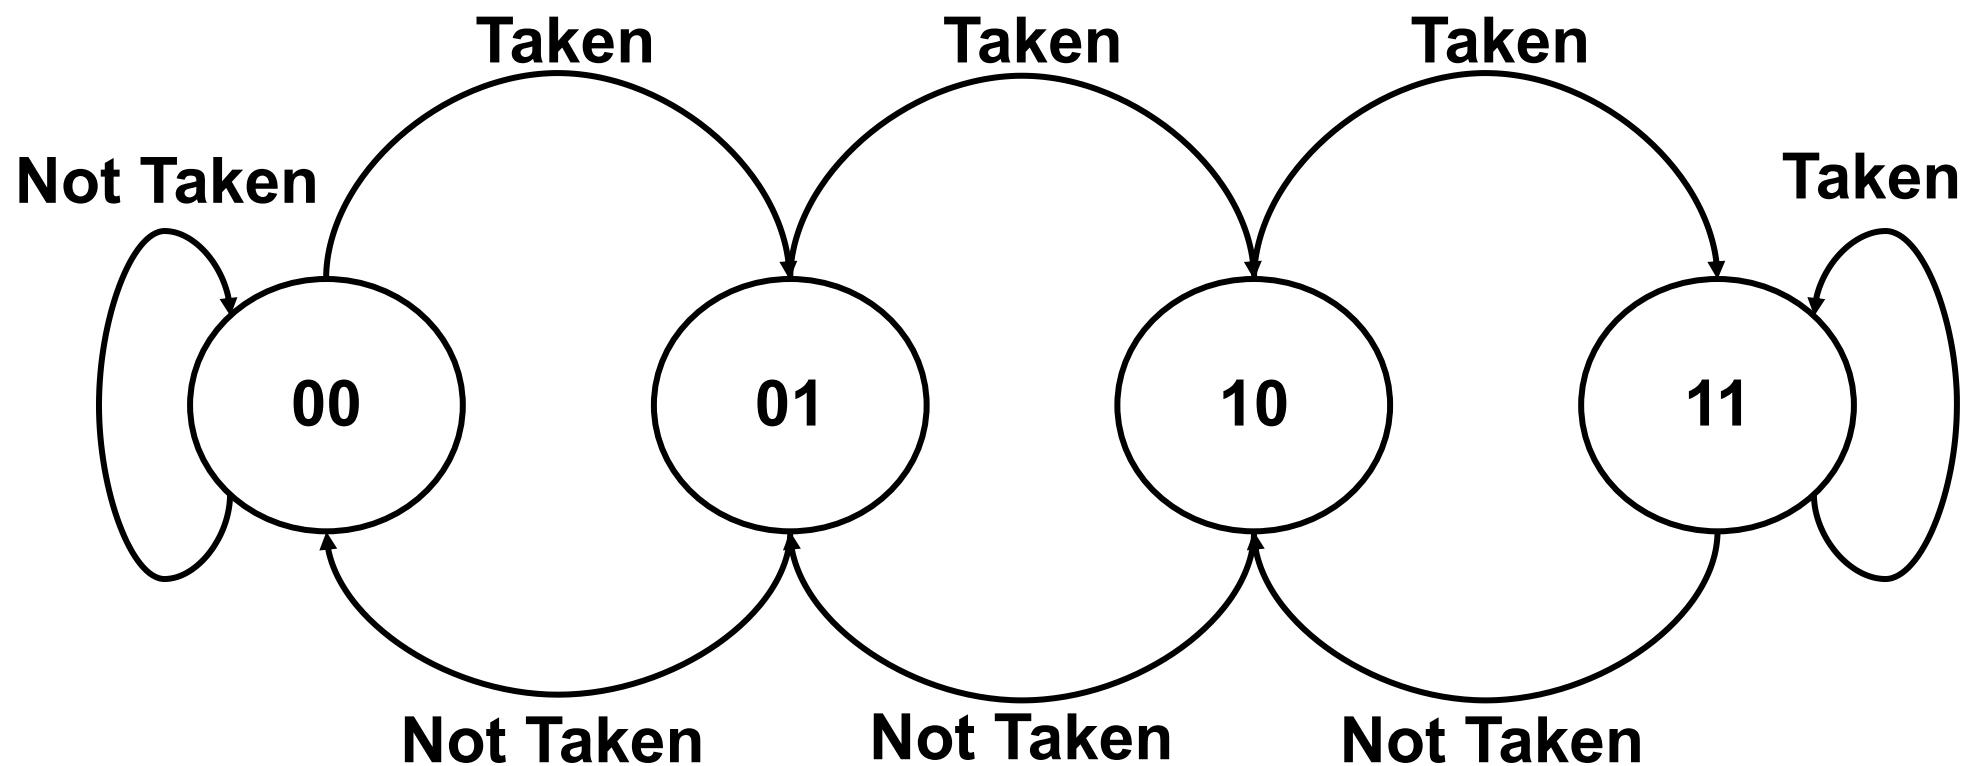
The counter is incremented each time the branch is taken, decremented each time it is not.

When the counter value is greater than one half of its maximum value, the branch is predicted as taken, when it is lower it is predicted as not taken.

Experiments have shown that there is little advantage in using $n > 2$.

This technique is also called bimodal branch prediction.

2-bit saturating counter



Performance impact of branches

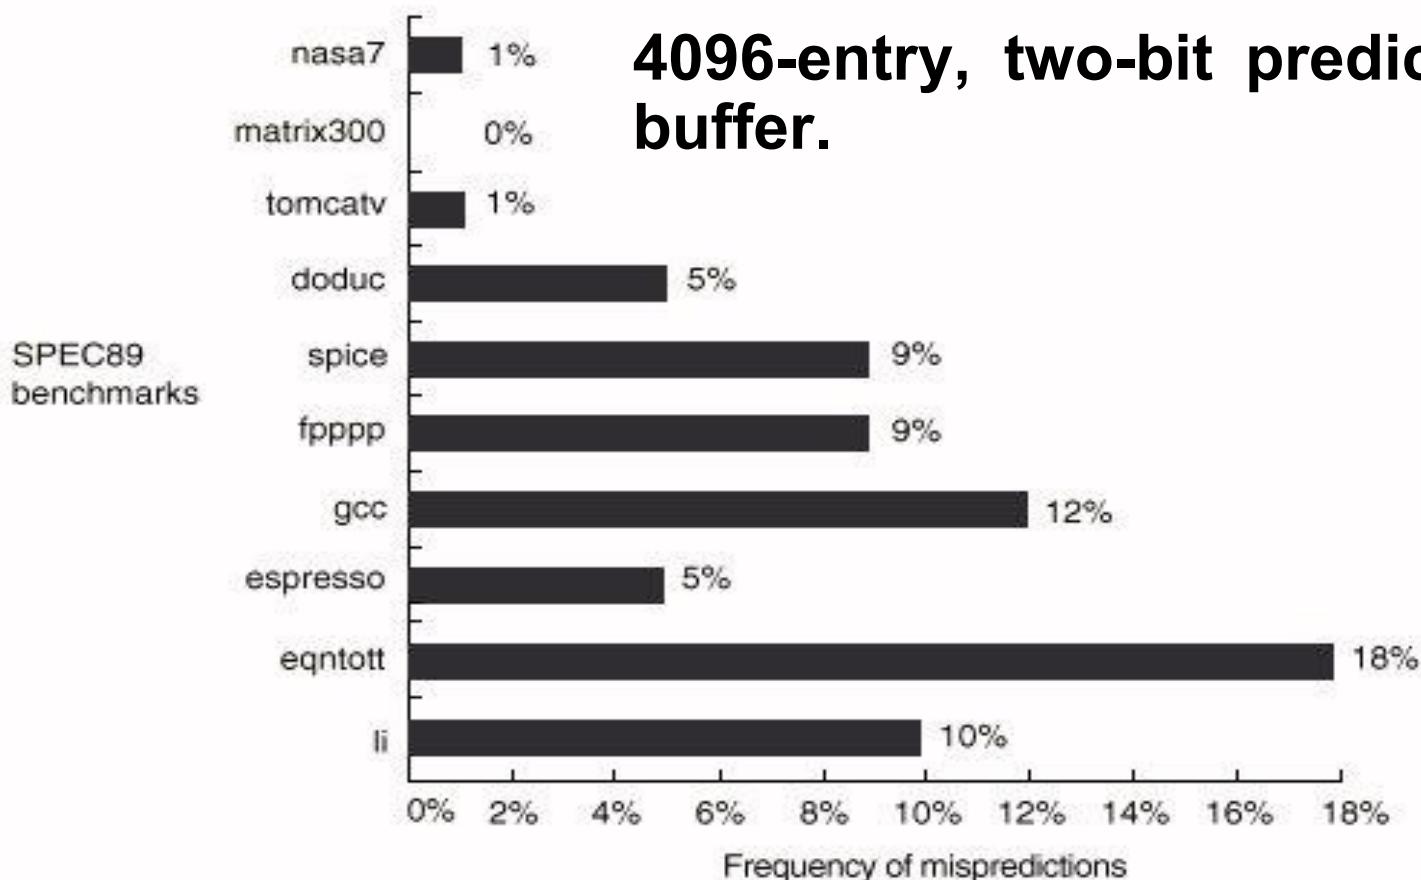
It depends on

- prediction accuracy
- branch cost (penalty for misprediction)
- branch frequency (lower for FP programs).

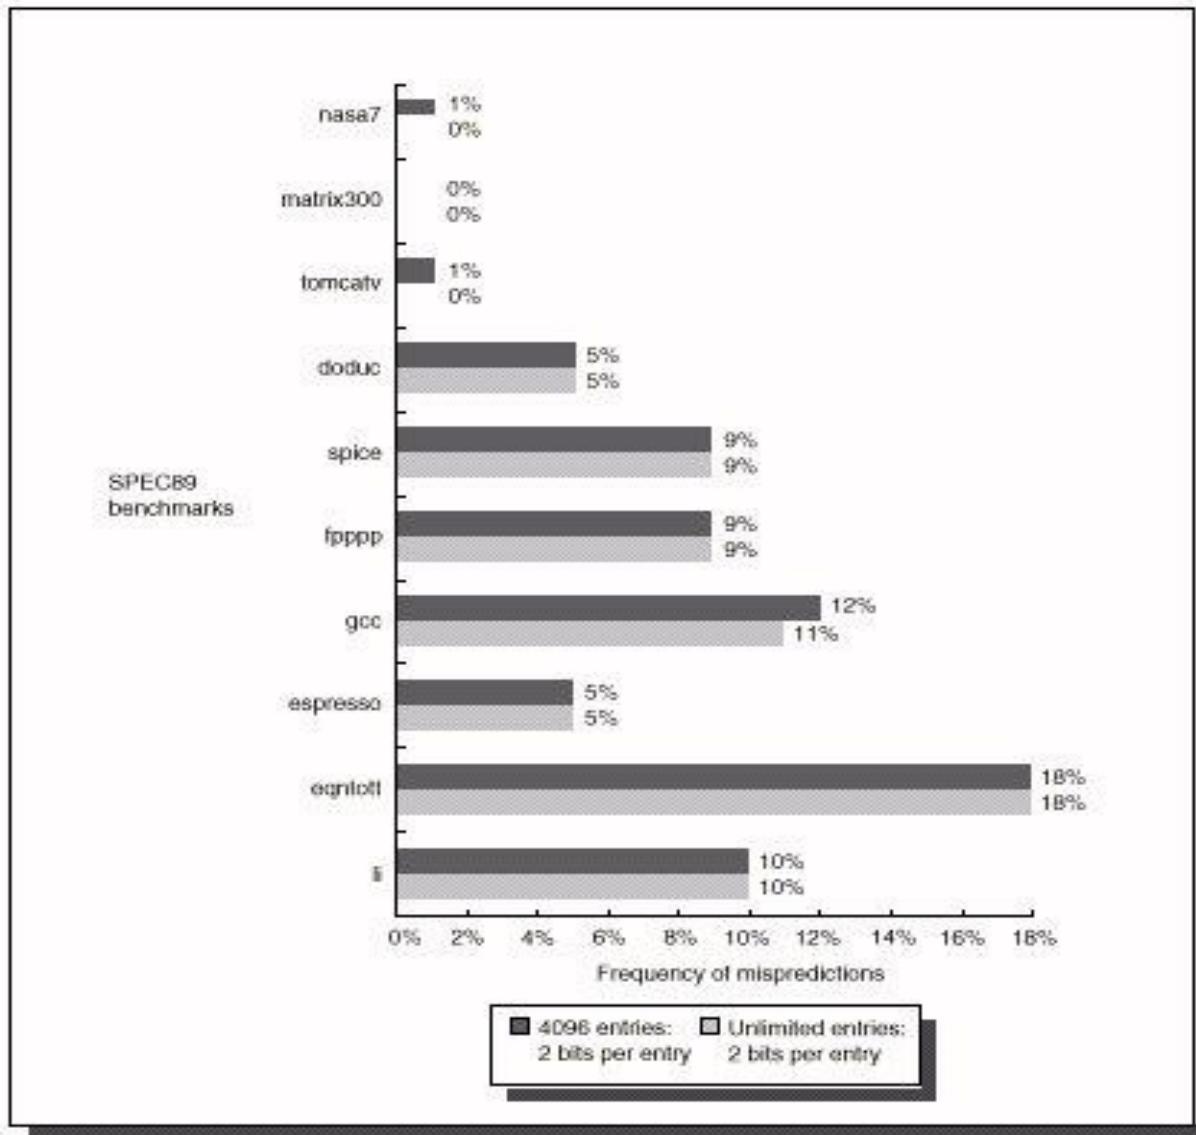
Prediction Accuracy for SPEC89

Assumption:

4096-entry, two-bit prediction buffer.



Dependence on Buffer Size



CORRELATING PREDICTORS

This approach (also called *two-level predictors*) is based on the dependencies between the results of last branches.

Example

```
if (aa==2)
    aa = 0;
if (bb==2)
    bb = 0;
if (aa !== bb)
{ }
```

The behavior of the last branch is strongly dependent on the result of the previous ones.

(m,n) predictors

They use the behavior of the last m branches to choose from 2^m branch predictors, each of which is a n -bit predictor.

The hardware required for implementing this scheme is very simple:

- The history of the most recent m branches is recorded in an m -bit shift register, where each bit records whether the branch was taken or not.
- The branch-prediction buffer is indexed using a concatenation of the low-order bits from the branch address with the m low-order history bits.

(1,1) predictor

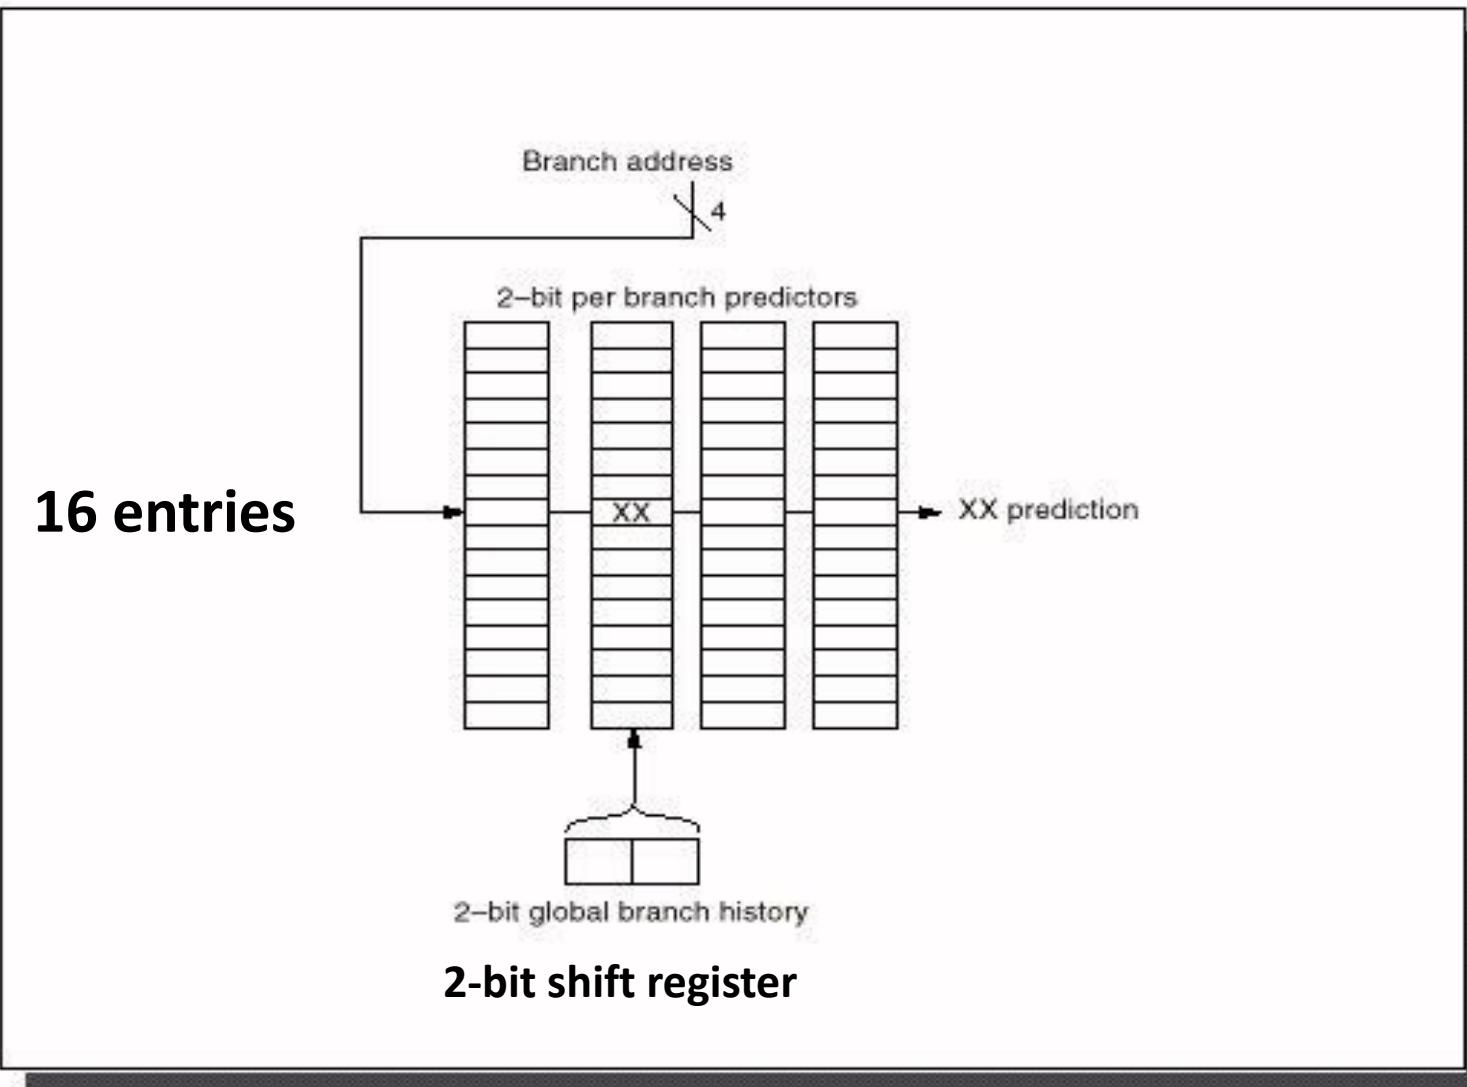
In this case: $m = 1$, $n = 1$.

Each branch is associated with 2^m (two) predictors of n (one) bits:

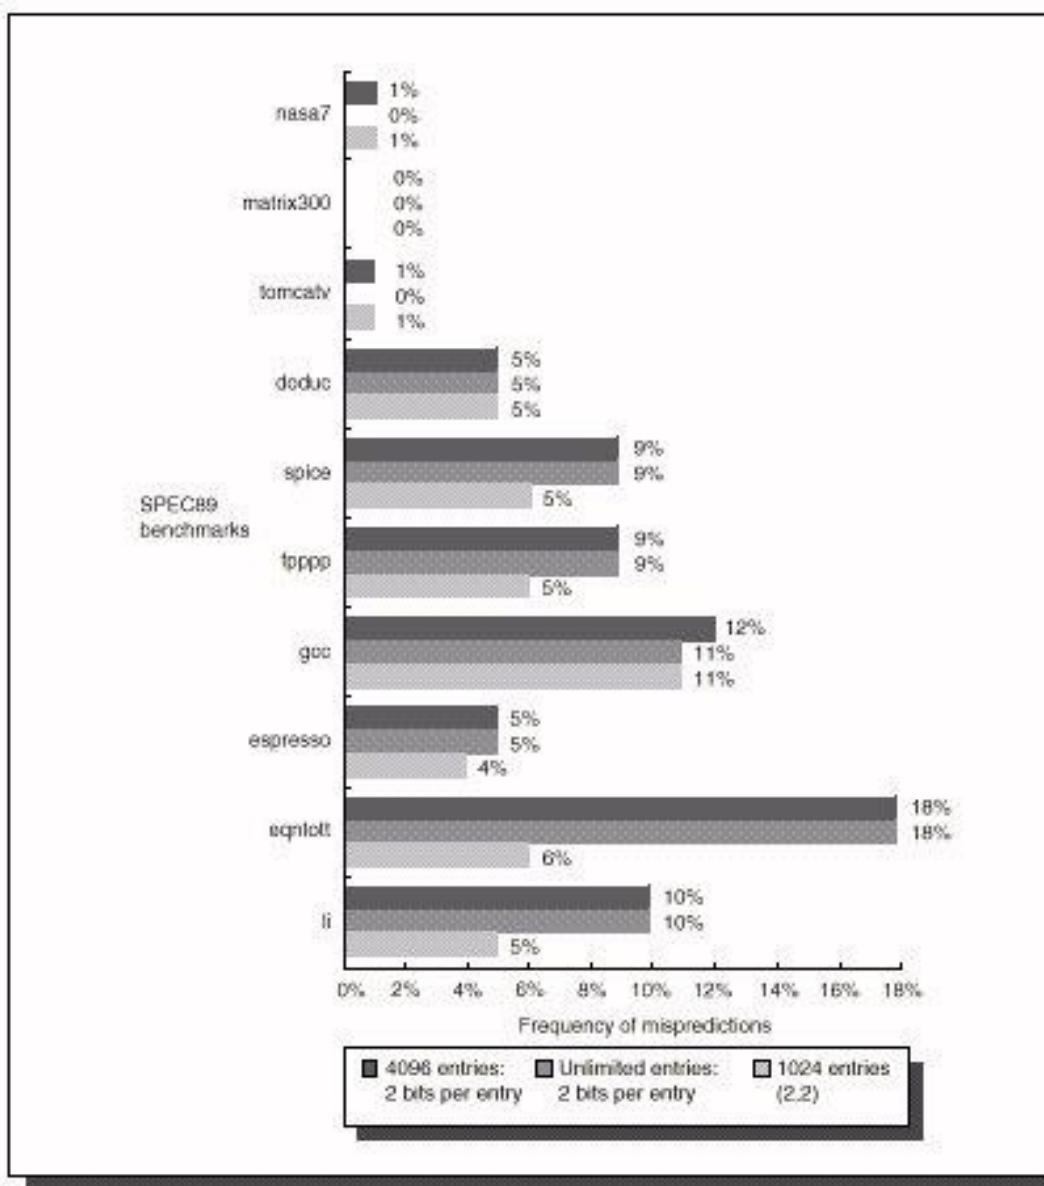
- one reporting the prediction in the case the previous branch was taken
- one with the prediction in the case the previous branch was not taken.

Warning: the two branches the prediction is based on can be different.

(2,2) predictor implementation



Performance comparison



Experimental Evidence

Prediction accuracy does not grow significantly by increasing

- the buffer size
- the number of bits per predictor.

BRANCH-TARGET BUFFER

Reducing the negative effects of control dependencies requires knowing as soon as possible

- whether the branch has to be taken or not
- the new value of the PC (if the branch is assumed to be taken).

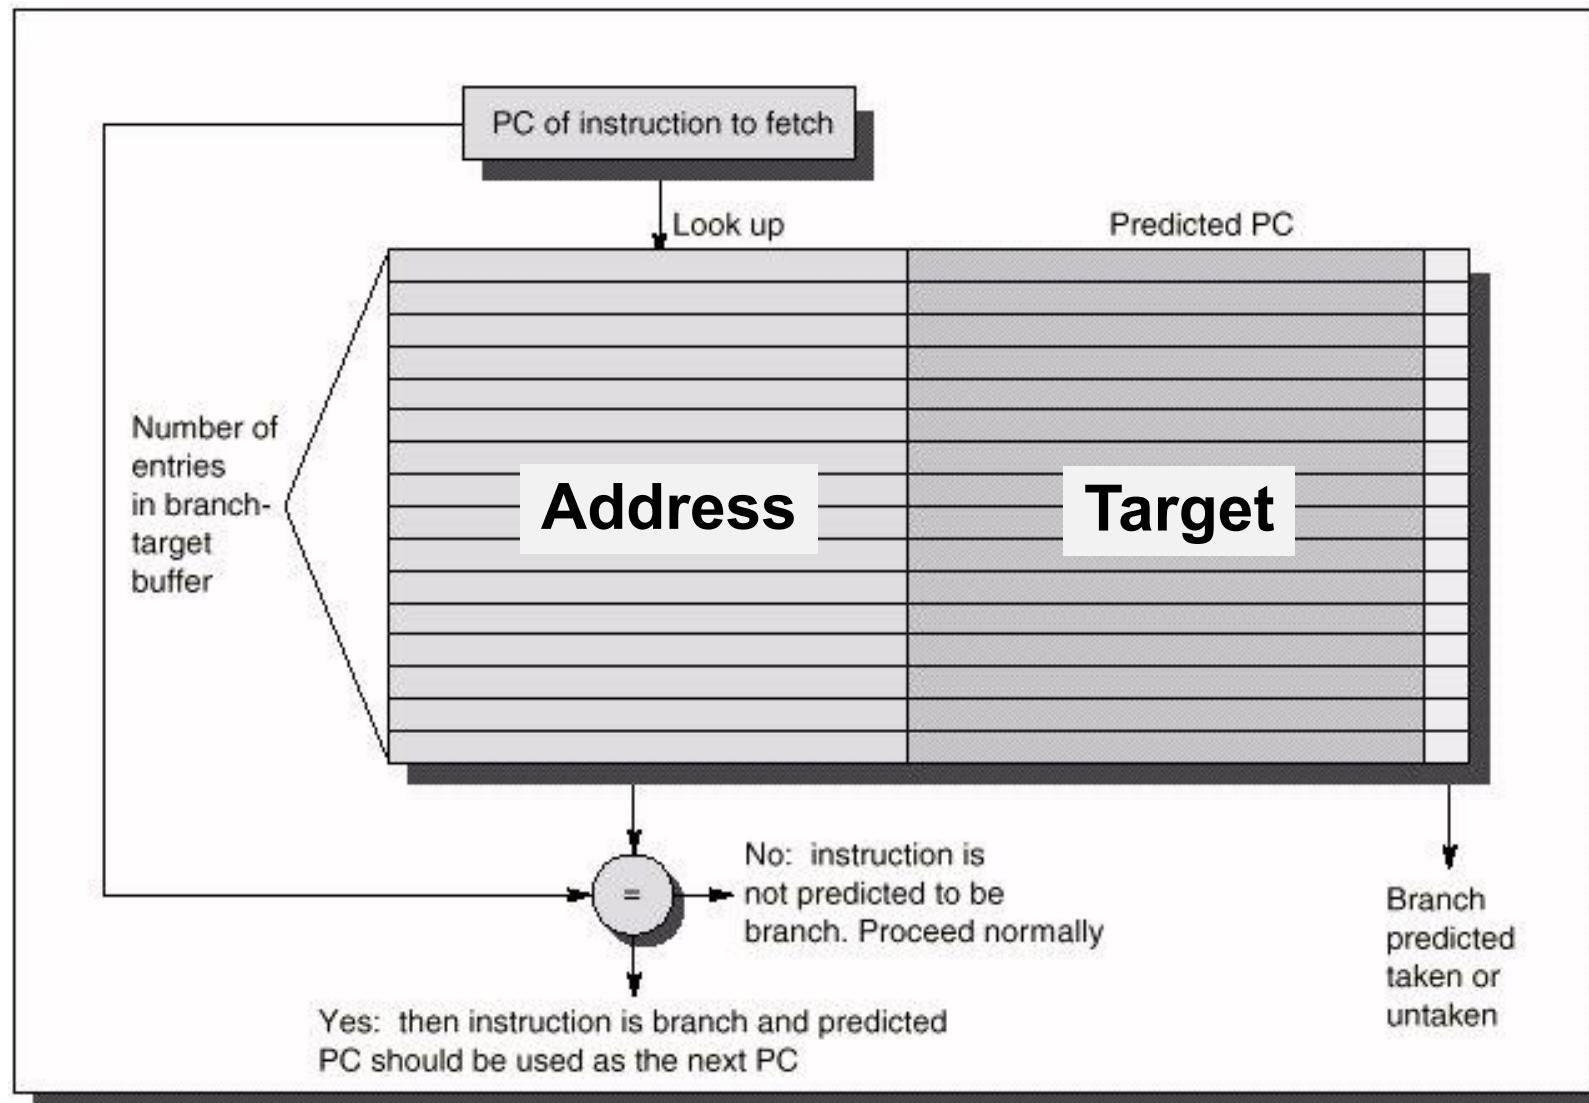
The later issue is faced by introducing a *branch-target buffer* (or *cache*).

Each entry of the branch-target buffer contains

- the **address** of the considered branch
- the **target** value to be loaded in the PC.

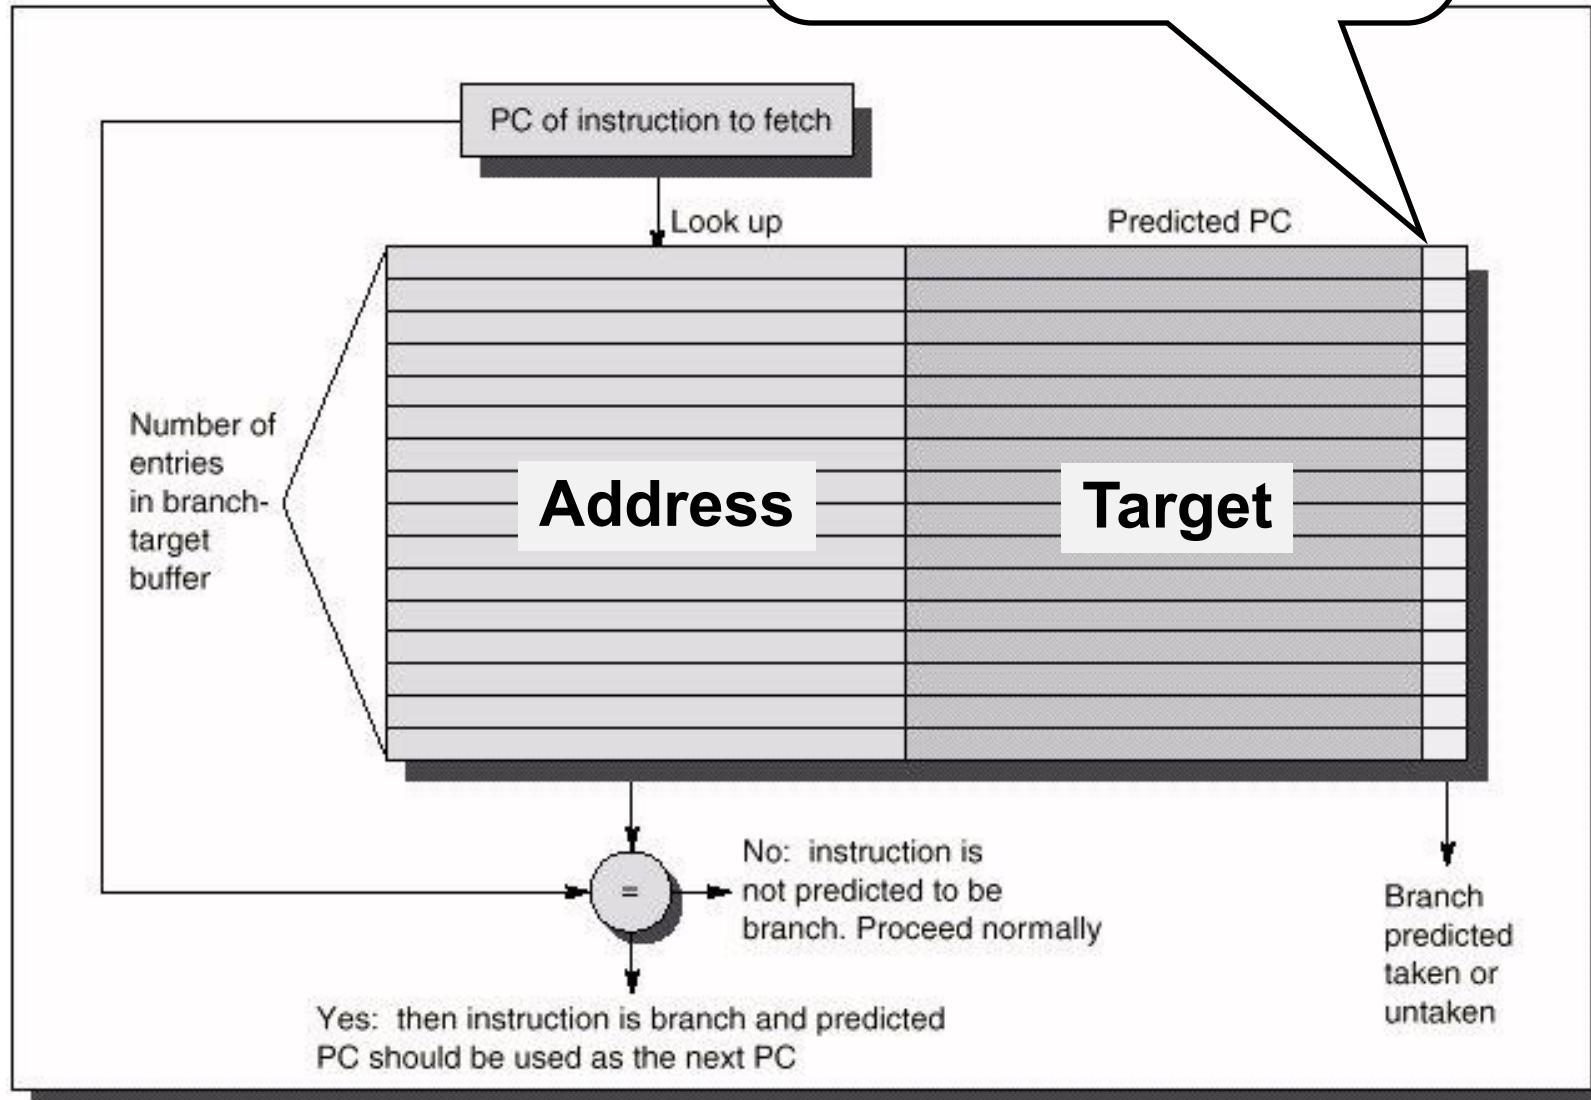
Using the branch-target buffer, the PC is loaded with the new value at the end of the IF stage, i.e., even before the branch instruction is decoded.

Branch-target Buffer: Architecture

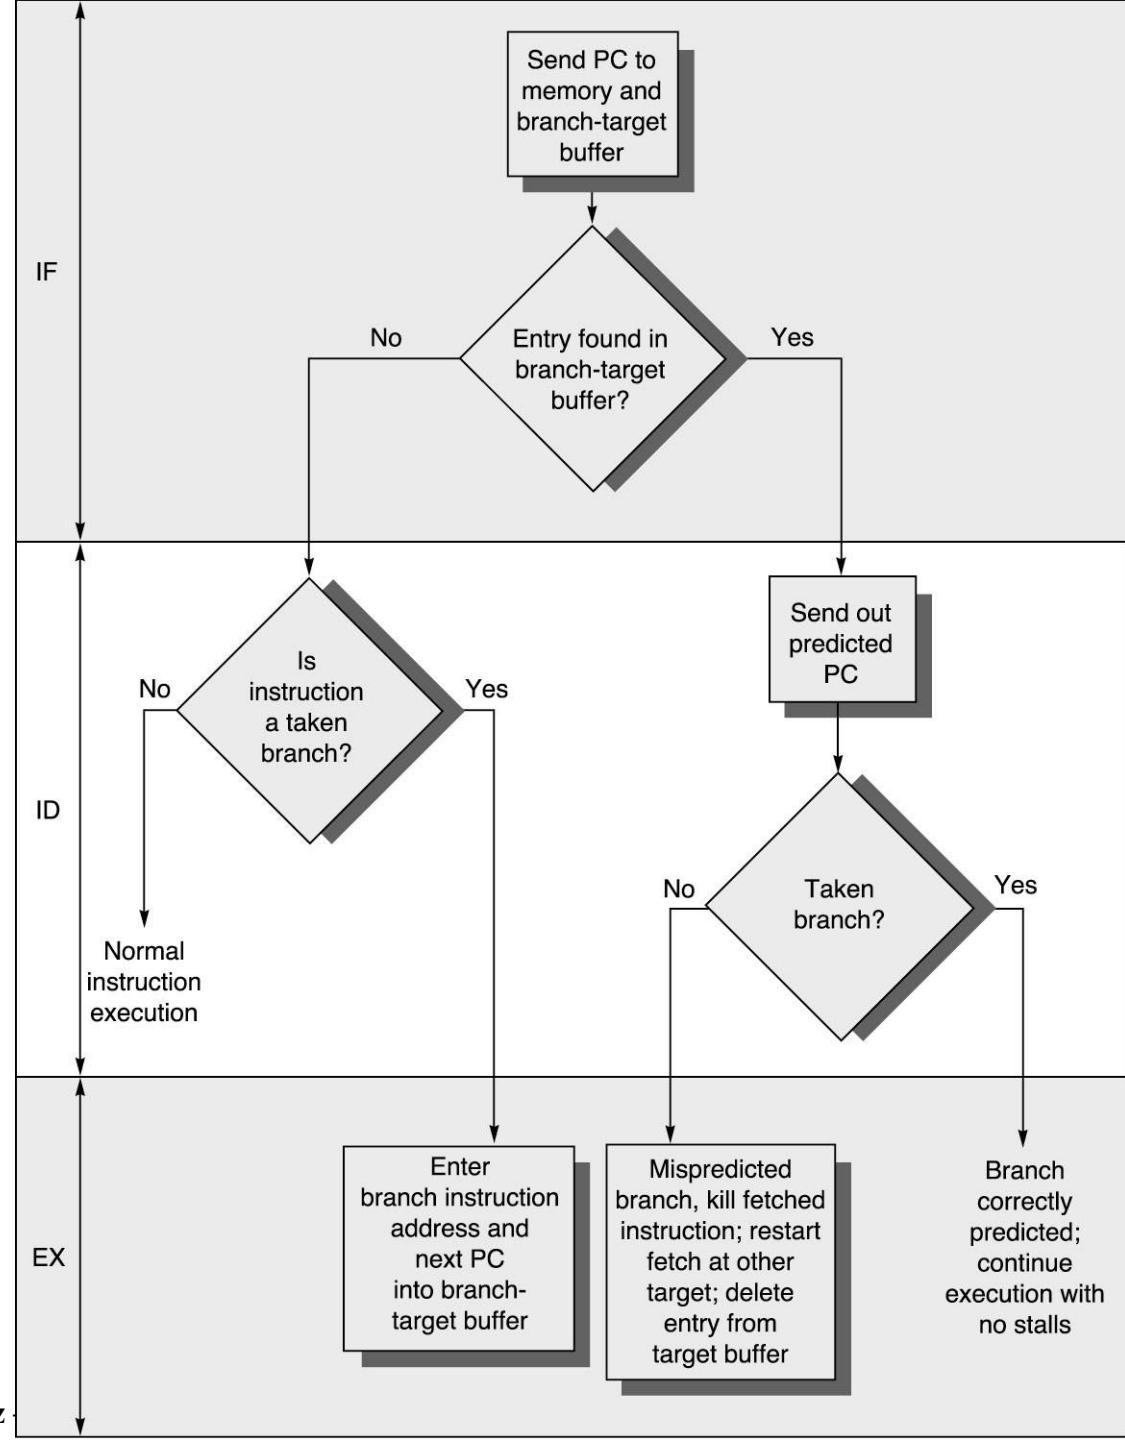


Branch-target Buffer

This field is not strictly necessary.



Branch-target Buffer: Behavior



Branch-target buffer: Advanced Issues

If a two-bit prediction strategy is adopted, it is possible to combine a branch-target buffer with a branch prediction buffer (i.e., branch history table).

Branch-target buffer: Performance Effects

Let assume the following penalty parameters

Instruction in buffer	Prediction	Actual branch	Penalty c.c.
Yes	taken	taken	0
Yes	taken	not taken	2
No		taken	2
No		not taken	0

Let also assume that

- the prediction accuracy is 90%
- the hit rate in the buffer is 90%
- taken branches are 60%.

Which is the total branch penalty?

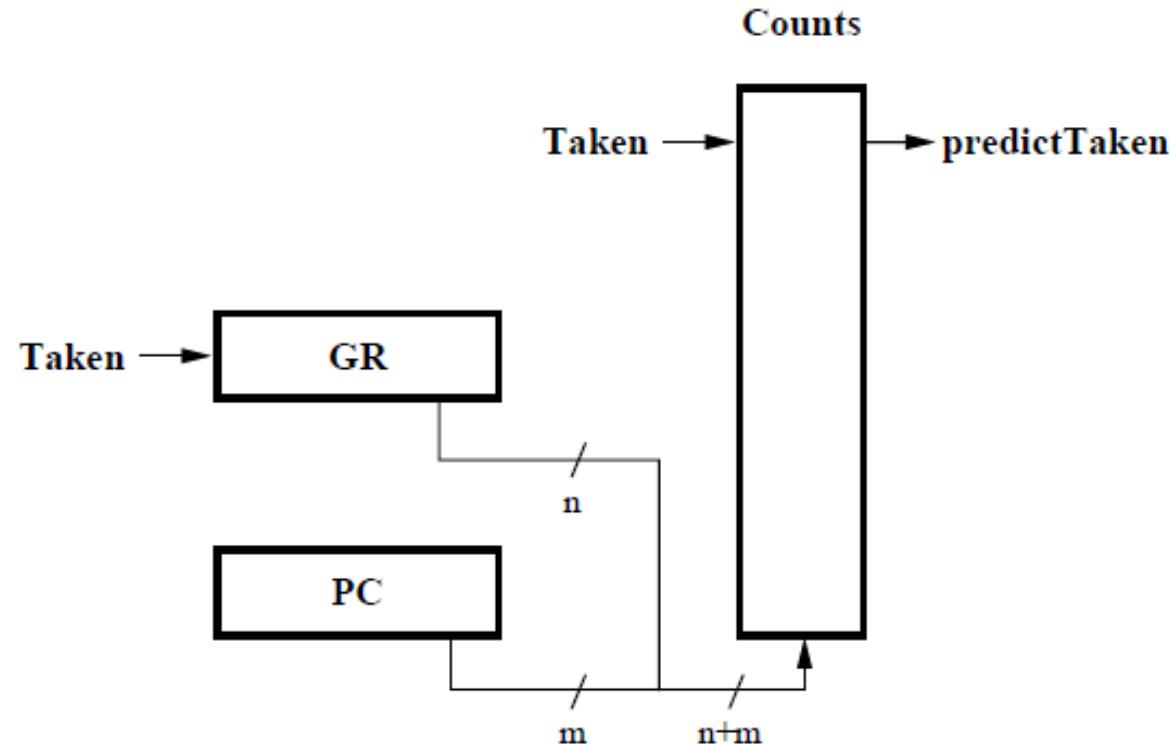
Solution

$$\begin{aligned}\text{Branch Penalty} &= \text{Hit branches \& Not Taken} + \\&\quad \text{Miss branches \& Taken} \\&= (\text{percent buffer hit rate} \times \\&\quad \text{percent incorrect predictions} \times 2) + \\&\quad ((1 - \text{percent buffer hit rate}) \times \\&\quad \text{taken branches} \times 2) \\&= (90\% \times 10\% \times 2) + (10\% \times 60\% \times 2) \\&= 0.18 + 0.12 = 0.30 \text{ c.c.}\end{aligned}$$

This figure should be compared with the 0.50 clock cycles per branch penalty existing with delayed branches.

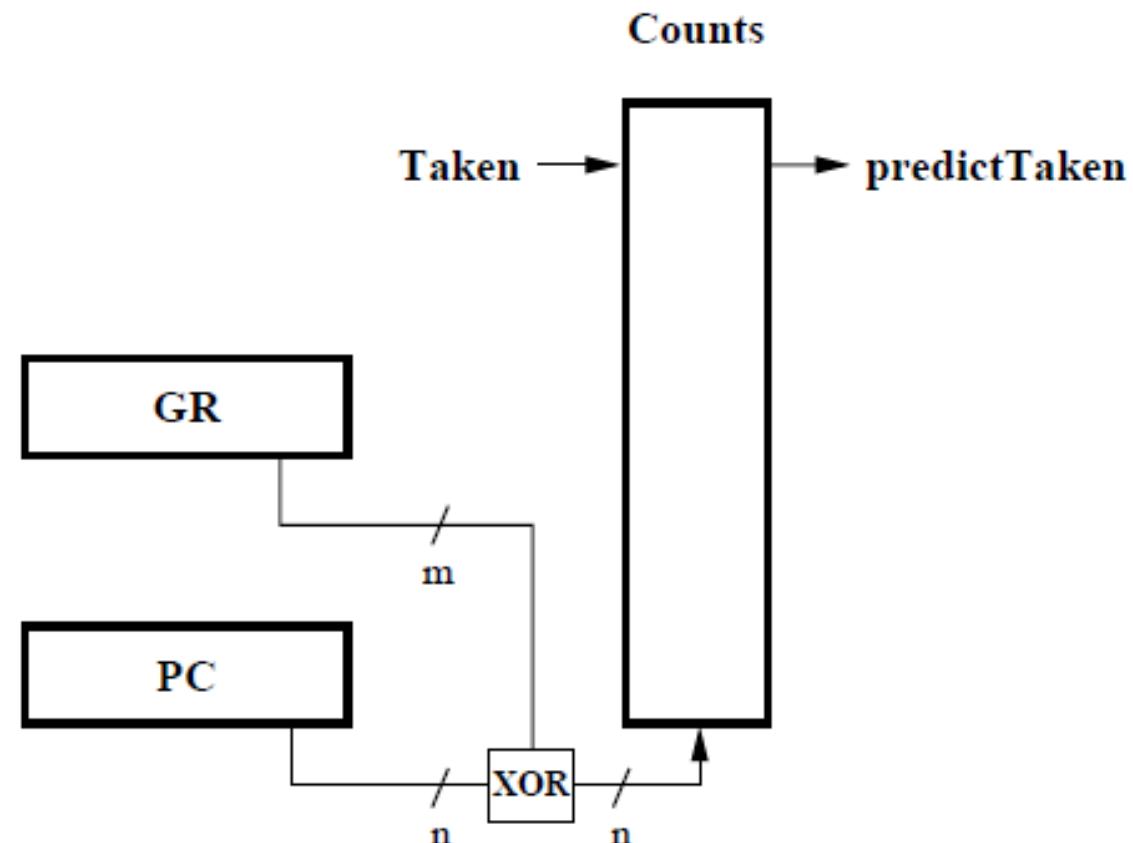
Global predictor with index selection (*gselect*)

A saturated-counter prediction table is accessed by concatenating the branches global history (GR) and the branch address (PC).

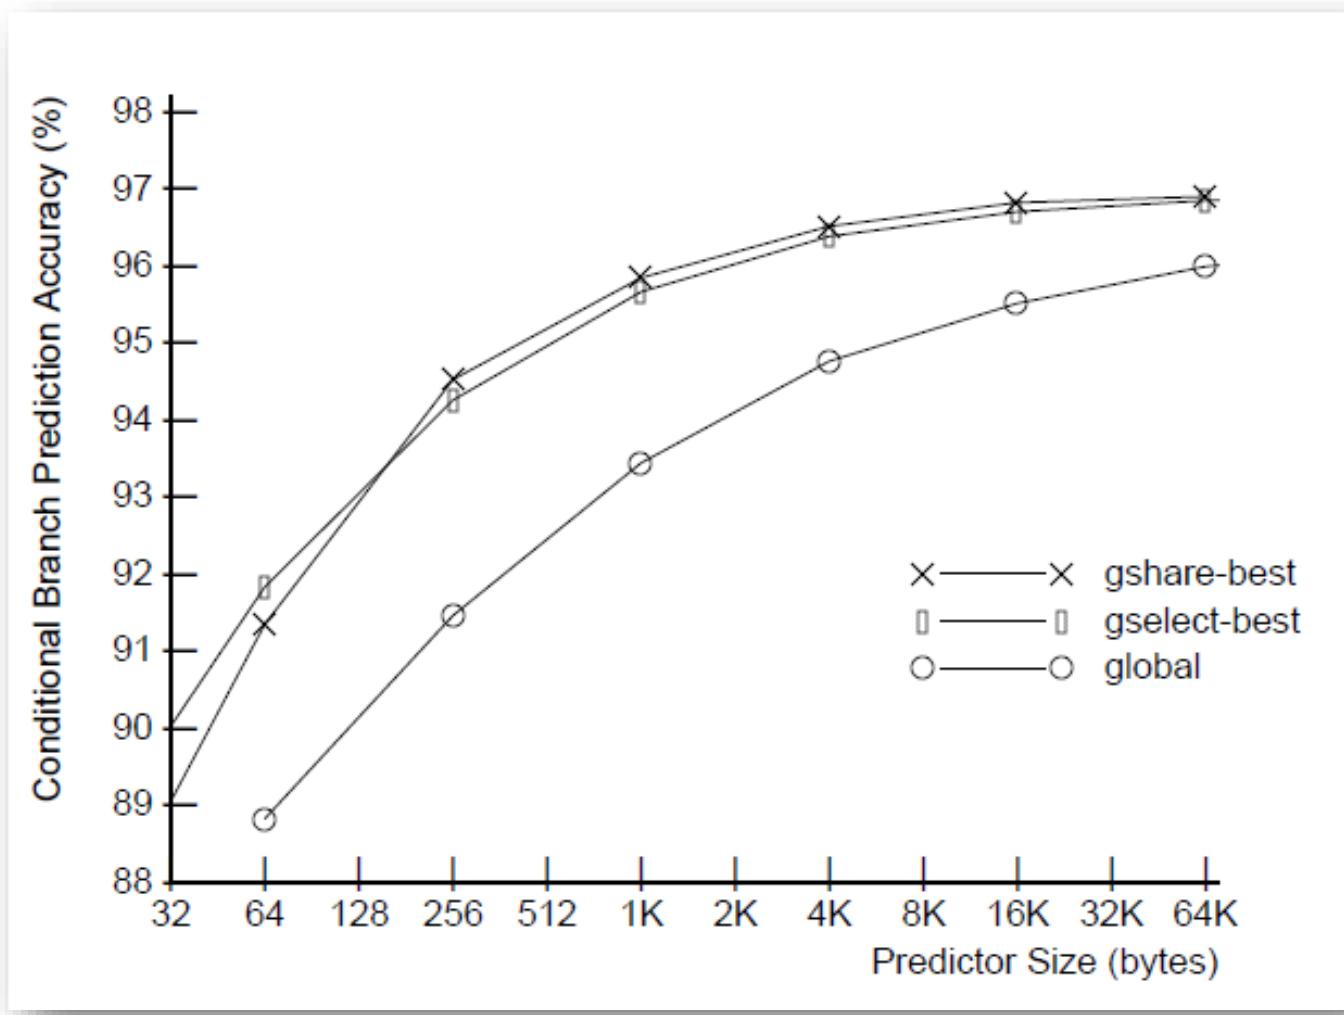


Global history with index sharing (*gshare*)

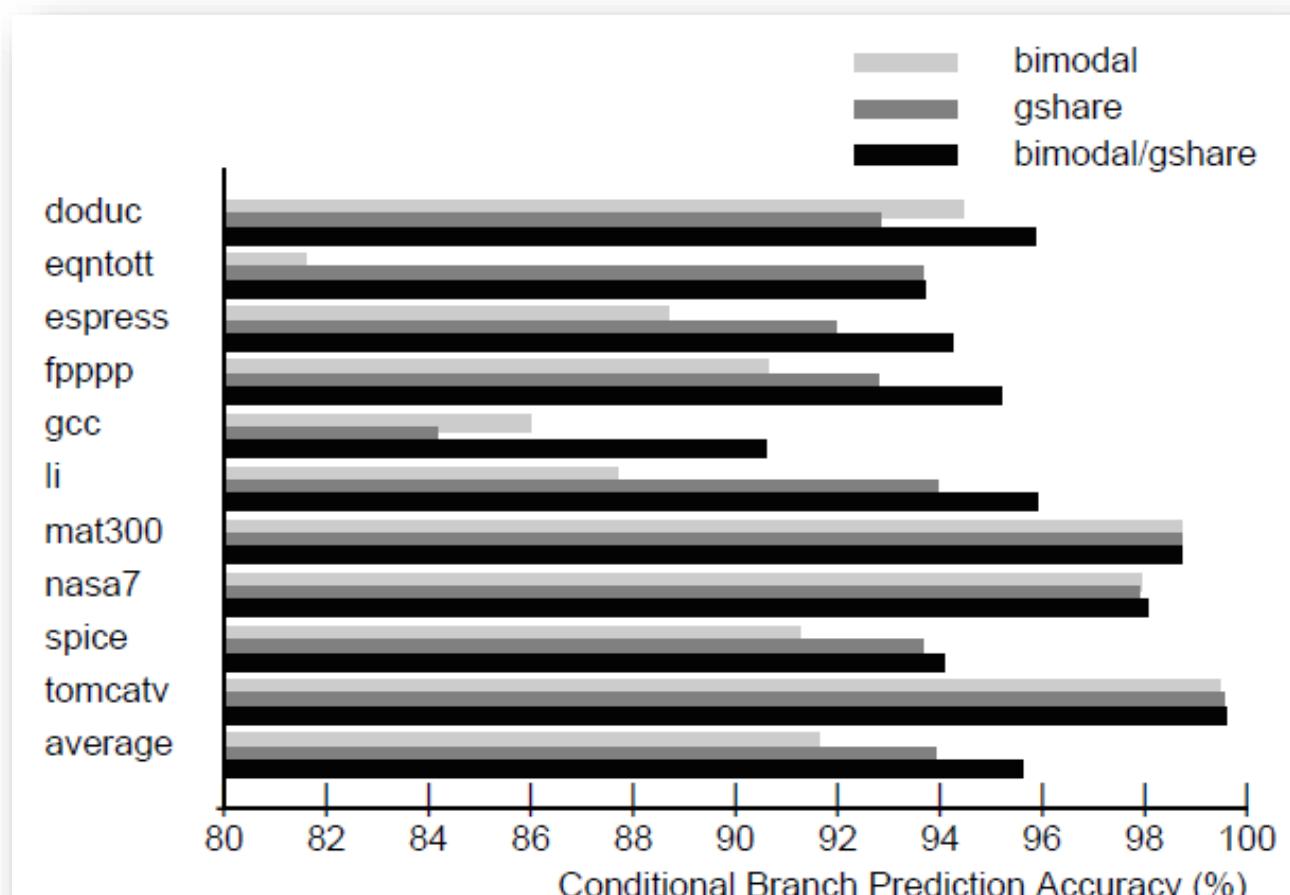
A saturated-counter prediction table is accessed by XORing the branches global history (GR) and the branch address (PC).



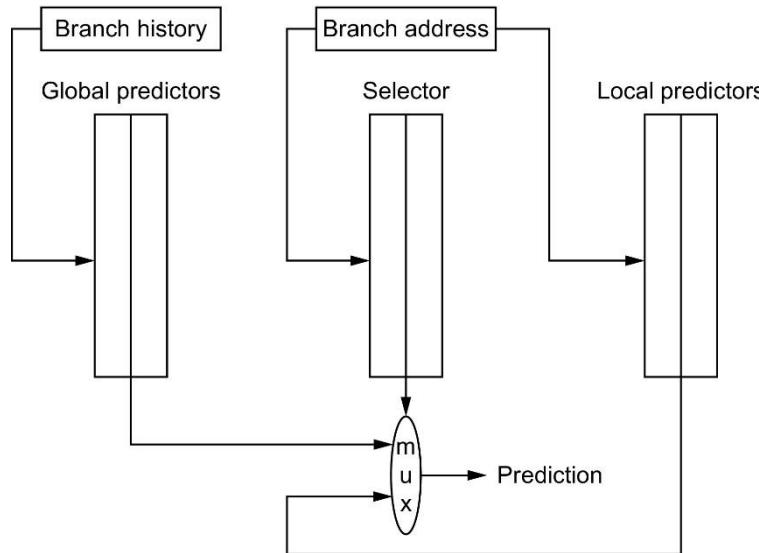
Performance comparison



Performance comparison



Tournament predictor



Global predictor: uses the history bits to index the predictors.

Local predictor: uses the address of the branch to index the predictors.

The tournament predictor combines *global* and *local* predictors, selected through a 2-bit saturated counter per branch to choose among the most effective of the two predictors.

Performance comparison

